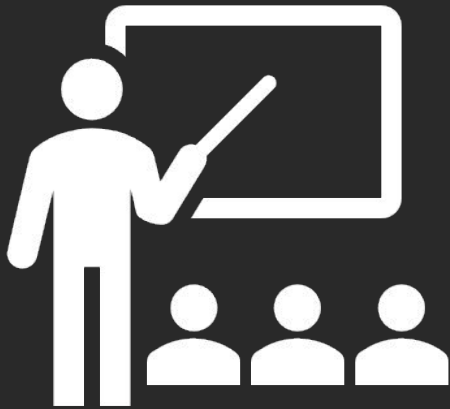


Containers

CS 106L Spring 2020 – Avery Wang and Anna Zeng
Stanford University

20 April 2020

Game Plan



- sequence containers
- container adaptors
- associative containers

Supplemental Material

- Regular slides: stuff we will go through lecture, absolutely necessary to understand future lectures.
- Supplemental material: important to know if you want to be considered "proficient" in C++ by the end of the class.
- Highly recommend you review the supplemental material.

Key questions we will answer today

- when should each container be used?
- how is the STL different from the Stanford libraries?
- what are some best practices and common pitfalls?
- homework: practice using the containers to solve problems.
- if time permits at the end, I'll do one problem

We'll assume you're familiar with

Stanford Library

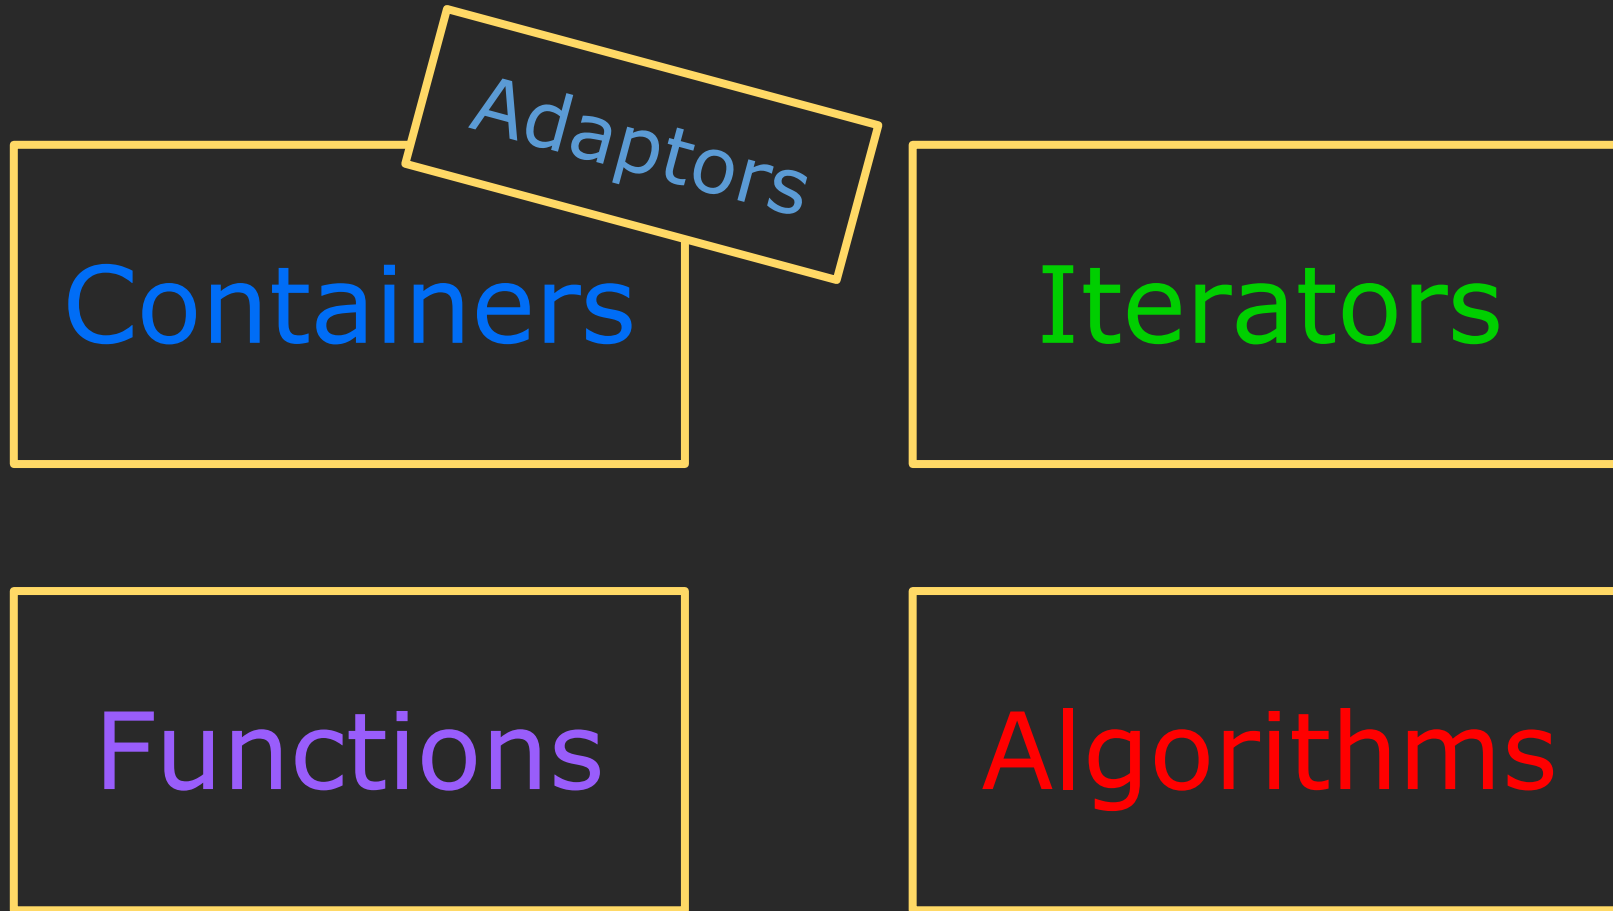
- `Vector<T>`
- `Stack<T>`, `Queue<T>`
- `Map<T>`, `Set<T>`

C++ details we will cover

Library

- `std::vector`
- `std::deque/list`
- `std::map/unordered_map`
- `std::set/unordered_set`
- `std::stack/queue`

Overview of STL



Sequence Containers

1. `std::vector` - Stanford vs. STL
2. `std::vector` - safety and efficiency
3. `std::deque` vs. `std::vector`

Stanford vs. STL vector: very similar!

```
1. // Stanford
2. Vector<char> vec{'a', 'b', 'c'};
3.
4. vec[0] = 'A';
5. cout << vec[vec.size()-1];
6.
7. for (int i = 0; i < vec.size(); ++i) {
8.     ++vec[i];
9. }
10.
11. for (auto& elem : vec) {
12.     --elem;
13. }
```

```
1. // STL
2. std::vector<char> vec{'a', 'b', 'c'};
3.
4. vec[0] = 'A';
5. cout << vec[vec.size()-1]; // or vec.back()
6.
7. for (size_t i = 0; i < vec.size(); ++i) {
8.     ++vec[i];
9. }
10.
11. for (auto& elem : vec) {
12.     --elem;
13. }
```

Answer on chat: what is different?

Stanford vs. STL vector: fairly close!

```
1. // Stanford
2. Vector<char> vec(3, 'c');
3. if (vec.isEmpty() || vec.size() == 3) {
4.     vec.clear();
5. }
6.
7. vec[-1] = 1; // BAD, error thrown!
8. Vector<char> vec2(3, 'c');
9. if (vec == vec2) {
10.     vec2.remove(vec2.size()-1);
11. }
```

```
1. // STL
2. std::vector<char> vec(3, 'c');
3. if (vec.empty() || vec.size() == 3) {
4.     vec.clear();
5. }
6.
7. vec[-1] = 1; // BAD, but no error thrown!
8. std::vector<char> vec2(3, 'c');
9. if (vec == vec2) {
10.     vec2.pop_back();
11. }
```

Answer on chat: what is different?

Stanford vs. STL vector: a summary

What you want to do	Stanford Vector<int>	std::vector<int>
Create an empty vector	Vector<int> v;	vector<int> v;
Create a vector with n copies of zero	Vector<int> v(n);	vector<int> v(n);
Create a vector with n copies of a value k	Vector<int> v(n, k);	vector<int> v(n, k);
Add k to the end of the vector	v.add(k);	v.push_back(k);
Clear vector	v.clear();	v.clear();
Get the element at index i (* does not bounds check!)	int k = v.get(i); int k = v[i];	int k = v.at(i); int k = v[i]; (*)
Check if the vector is empty	if (v.isEmpty()) ...	if (v.empty()) ...
Replace the element at index i (* does not bounds check!)	v.get(i) = k; v[i] = k;	v.at(i) = k; v[i] = k; (*)

Stanford vs. STL vector: a summary

What you want to do	Stanford Vector<int>	std::vector<int>
Create an empty vector	<code>Vector<int> v;</code>	<code>vector<int> v;</code>
Create a vector with n copies of zero	<code>Vector<int> v(n);</code>	<code>vector<int> v(n);</code>
Create a vector with n copies of a value k	<code>Vector<int> v(n, k);</code>	<code>vector<int> v(n, k);</code>
Add k to the end of the vector	<code>v.add(k);</code>	<code>v.push_back(k);</code>
Clear vector	<code>v.clear();</code>	<code>v.clear();</code>
Get the element at index i (* does not bounds check!)	<code>int k = v.get(i);</code> <code>int k = v[i];</code>	<code>int k = v.at(i);</code> <code>int k = v[i]; (*)</code>
Check if the vector is empty	<code>if (v.isEmpty()) ...</code>	<code>if (v.empty()) ...</code>
Replace the element at index i (* does not bounds check!)	<code>v.get(i) = k;</code> <code>v[i] = k;</code>	<code>v.at(i) = k;</code> <code>v[i] = k; (*)</code>

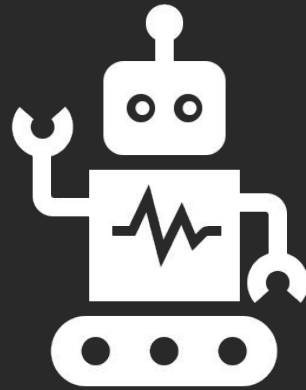
What is missing on this chart?

What you want to do	Stanford Vector<int>	std::vector<int>
Create an empty vector	<code>Vector<int> v;</code>	<code>vector<int> v;</code>
Create a vector with n copies of zero	<code>Vector<int> v(n);</code>	<code>vector<int> v(n);</code>
Create a vector with n copies of a value k	<code>Vector<int> v(n, k);</code>	<code>vector<int> v(n, k);</code>
Add k to the end of the vector	<code>v.add(k);</code>	<code>v.push_back(k);</code>
Clear vector	<code>v.clear();</code>	<code>v.clear();</code>
Get the element at index i (* does not bounds check!)	<code>int k = v.get(i);</code> <code>int k = v[i];</code>	<code>int k = v.at(i);</code> <code>int k = v[i]; (*)</code>
Check if the vector is empty	<code>if (v.isEmpty()) ...</code>	<code>if (v.empty()) ...</code>
Replace the element at index i (* does not bounds check!)	<code>v.get(i) = k;</code> <code>v[i] = k;</code>	<code>v.at(i) = k;</code> <code>v[i] = k; (*)</code>

Stanford vs. STL vector: a summary

What you want to do	Stanford Vector <int>	std:: vector <int>
Add j to the front of the vector	<code>v.insert(0, k);</code>	<code>v.insert(v.begin(), k);</code>
Insert k at some index i	<code>v.insert(i, k);</code>	<code>v.insert(v.begin()+i, k);</code>
Remove the element at index i	<code>v.remove(i);</code>	<code>v.erase(v.begin()+i);</code>
Get the sublist in indices [i, j)	<code>v.subList(i, j);</code>	<code>vector<int> c (v.begin()+i, v.begin()+j);</code>
Create a vector that is two vectors appended together.	<code>vector<int> v = v1 + v2;</code>	<code>// kinda complicated</code>

This'll require understanding iterators.
Next lecture!



Questions

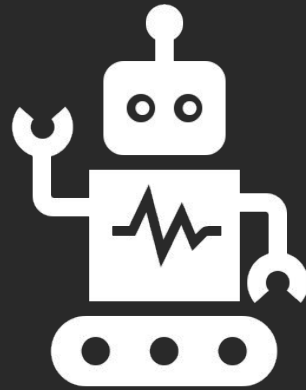
Answer 2 questions.

operator[] does not perform bounds checking

You'll see operator[] used, but rarely will you see at().

This follows the C++ philosophy to never waste time.

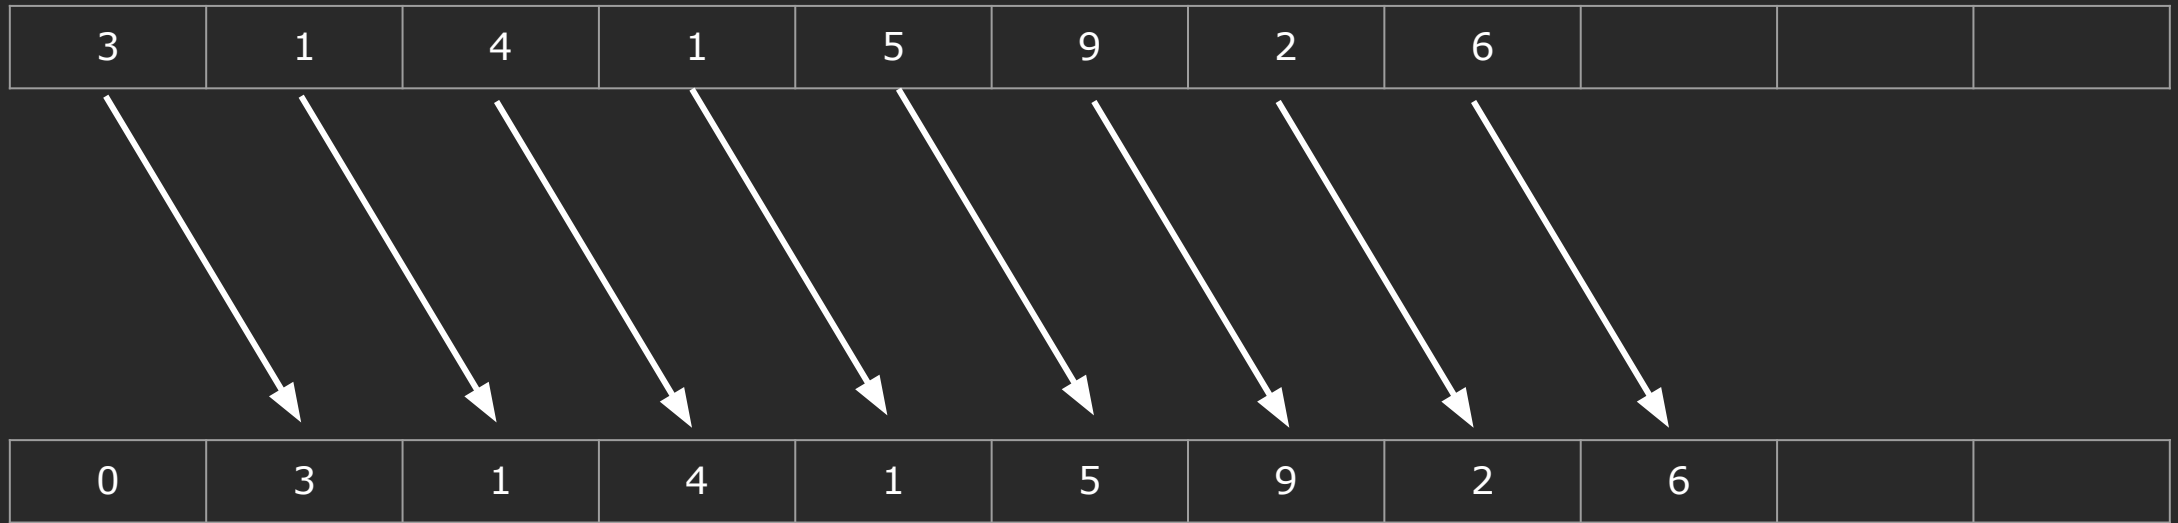
```
1. std::vector<int> vec{5, 6};    // {5, 6}
2. vec[1] = 3;                  // {5, 3}
3. vec[2] = 4;                  // undefined behavior
```

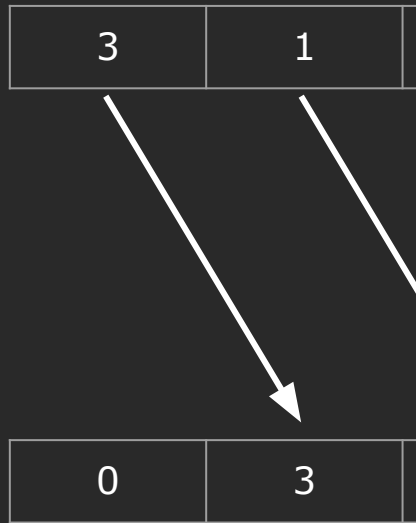
Example

Front insertion speed.

vector does not have a `push_front` function



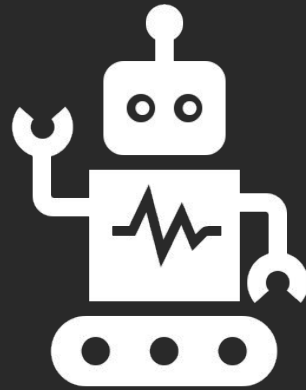
vector does not have a push_front function



vector does not have a push_front function

```
1. v.push_front(0); // not a real function
```

Recurring C++ pattern: don't provide functions which might be mistaken to be efficient when it's not.



Questions

Answer 2 questions.

What if you really wanted fast insertion to the front?

std::deque provides fast insertion anywhere

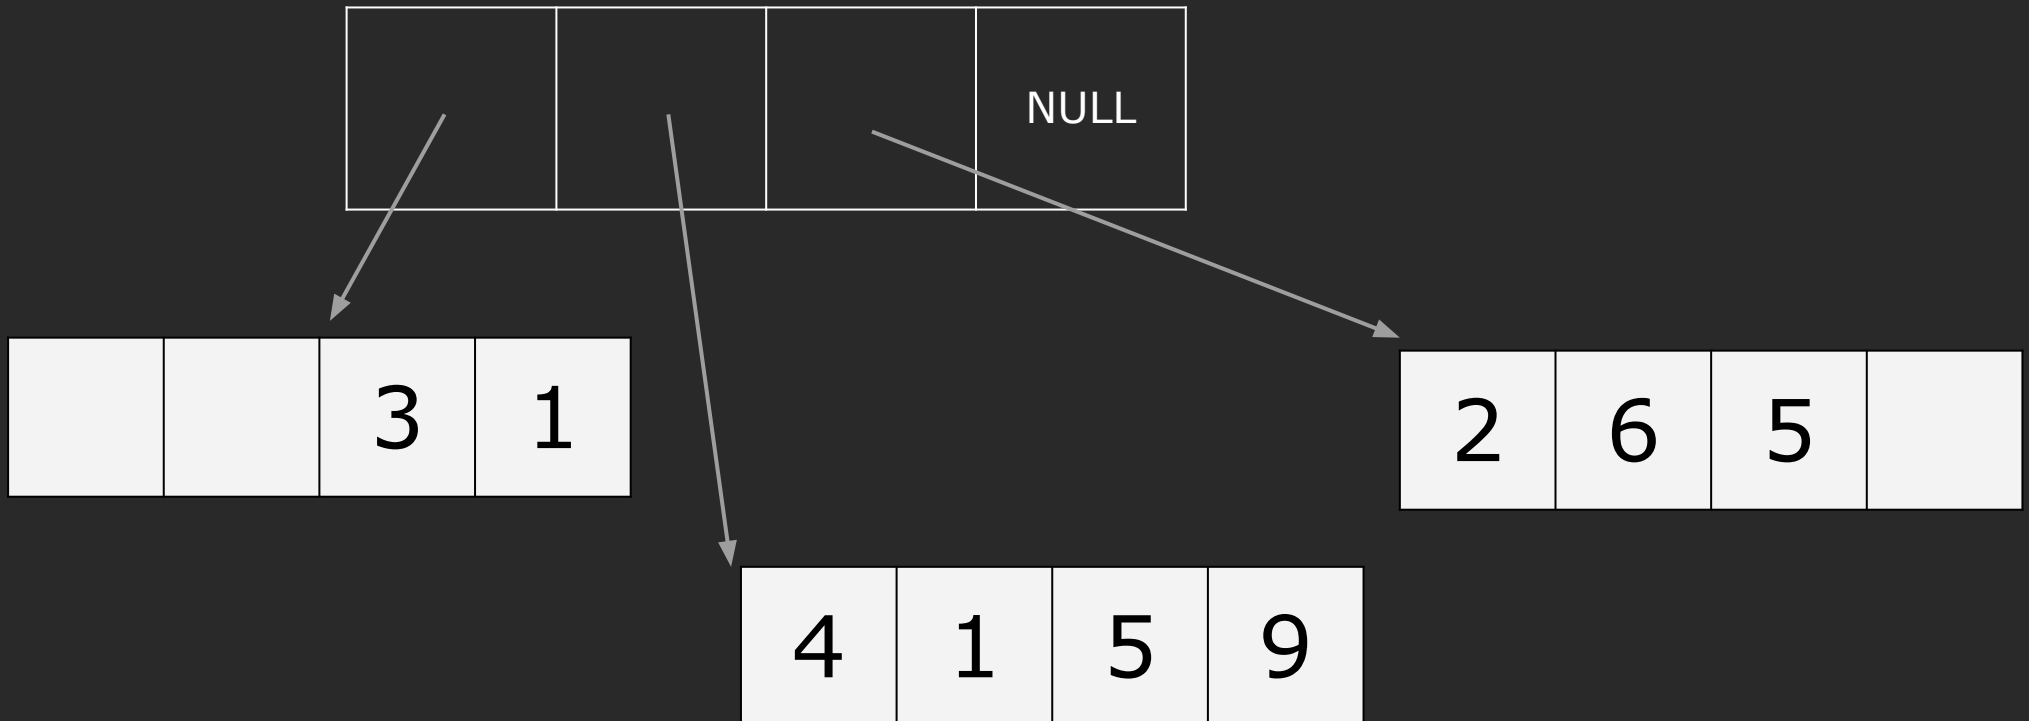
std::deque has the exact same functions as std::vector
but also has push_front and pop_front.

```
1. std::deque<int> deq{5, 6};           // {5, 6}
2. deq.push_front(3);                   // {3, 5, 6}
3. deq.pop_back(4);                     // {3, 5}
4. deq[1] = -2;                         // {3, -2}
```

How is a deque implemented?

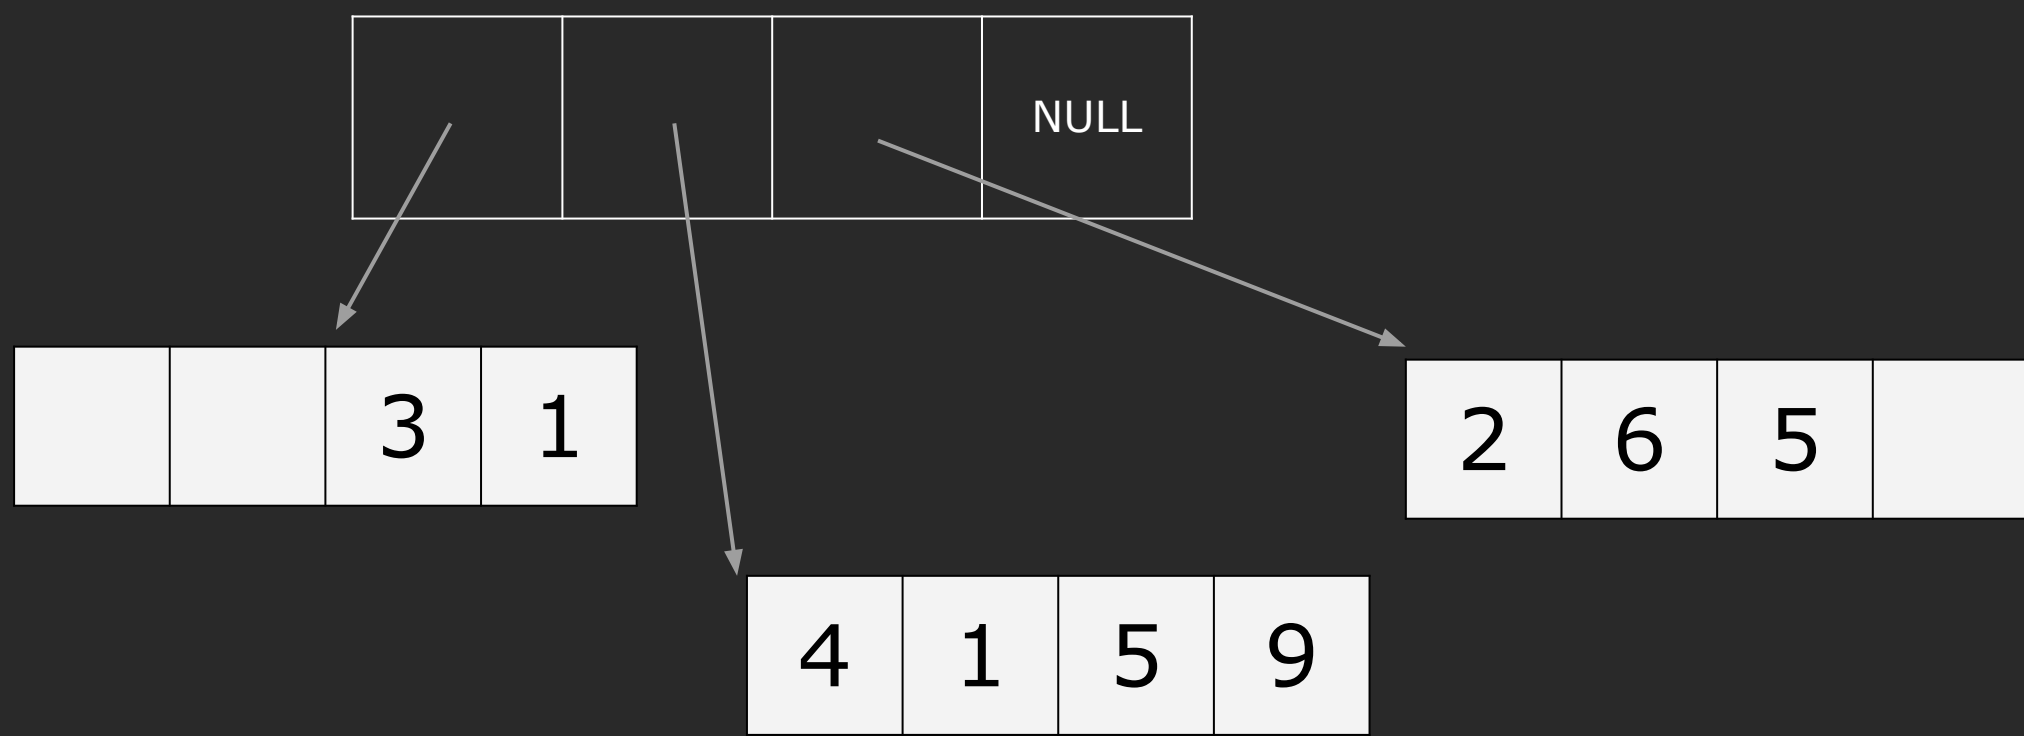
How does `std::deque<T>` work?

There is no single specific implementation of a deque, but one common one might look like this:



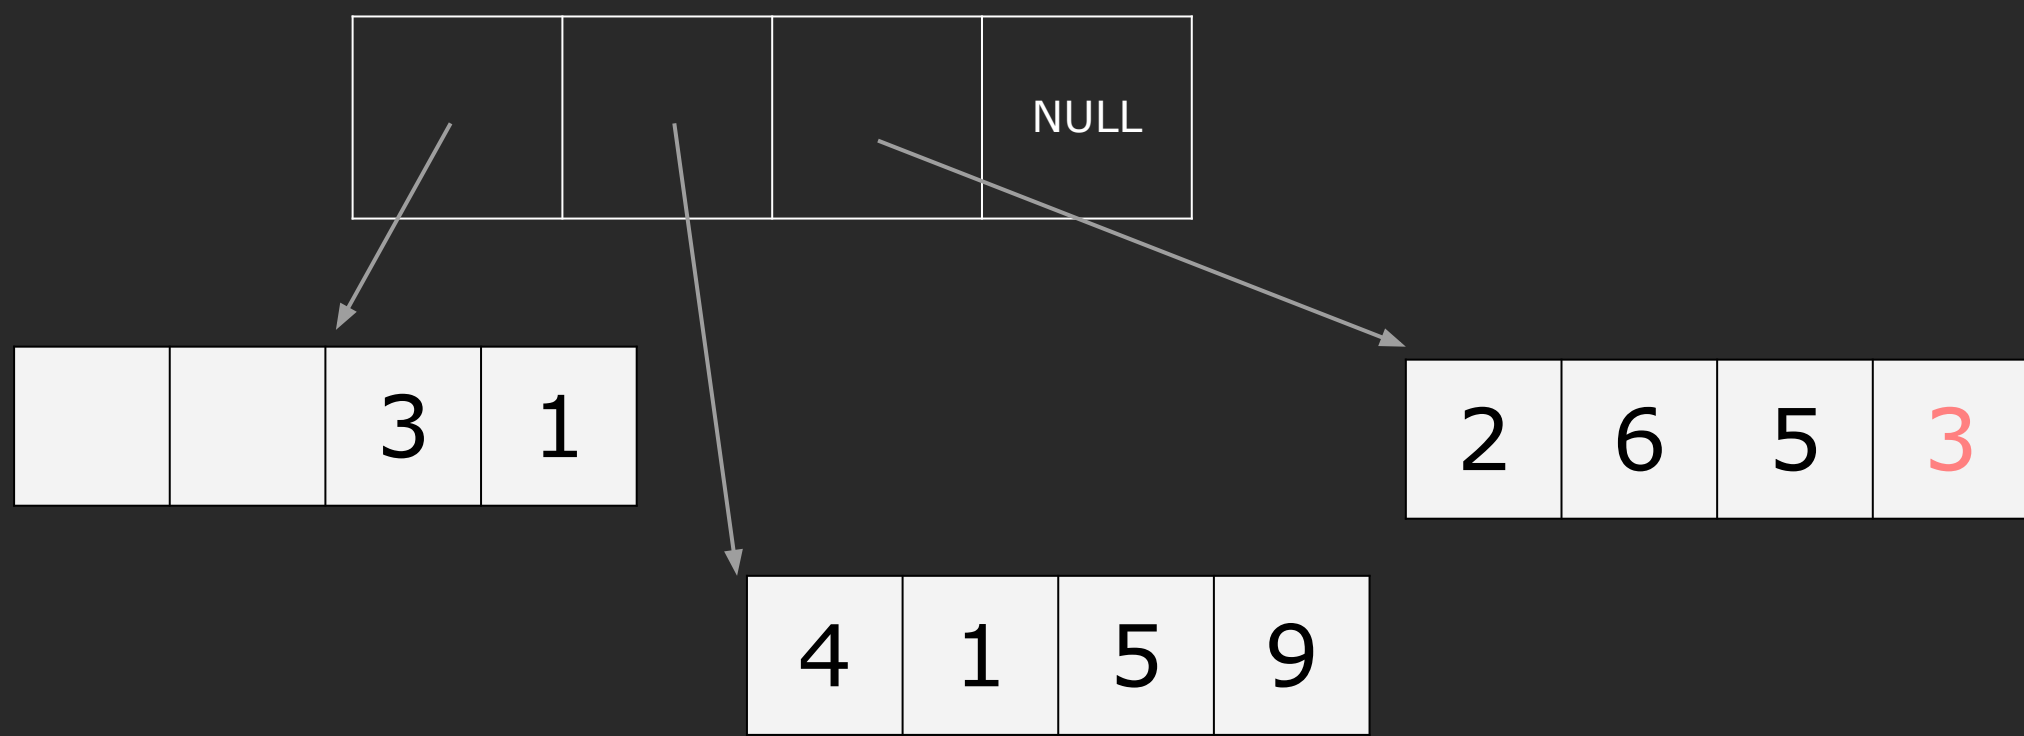
How does `std::deque<T>` work?

How would you do `push_back(3)`?



How does `std::deque<T>` work?

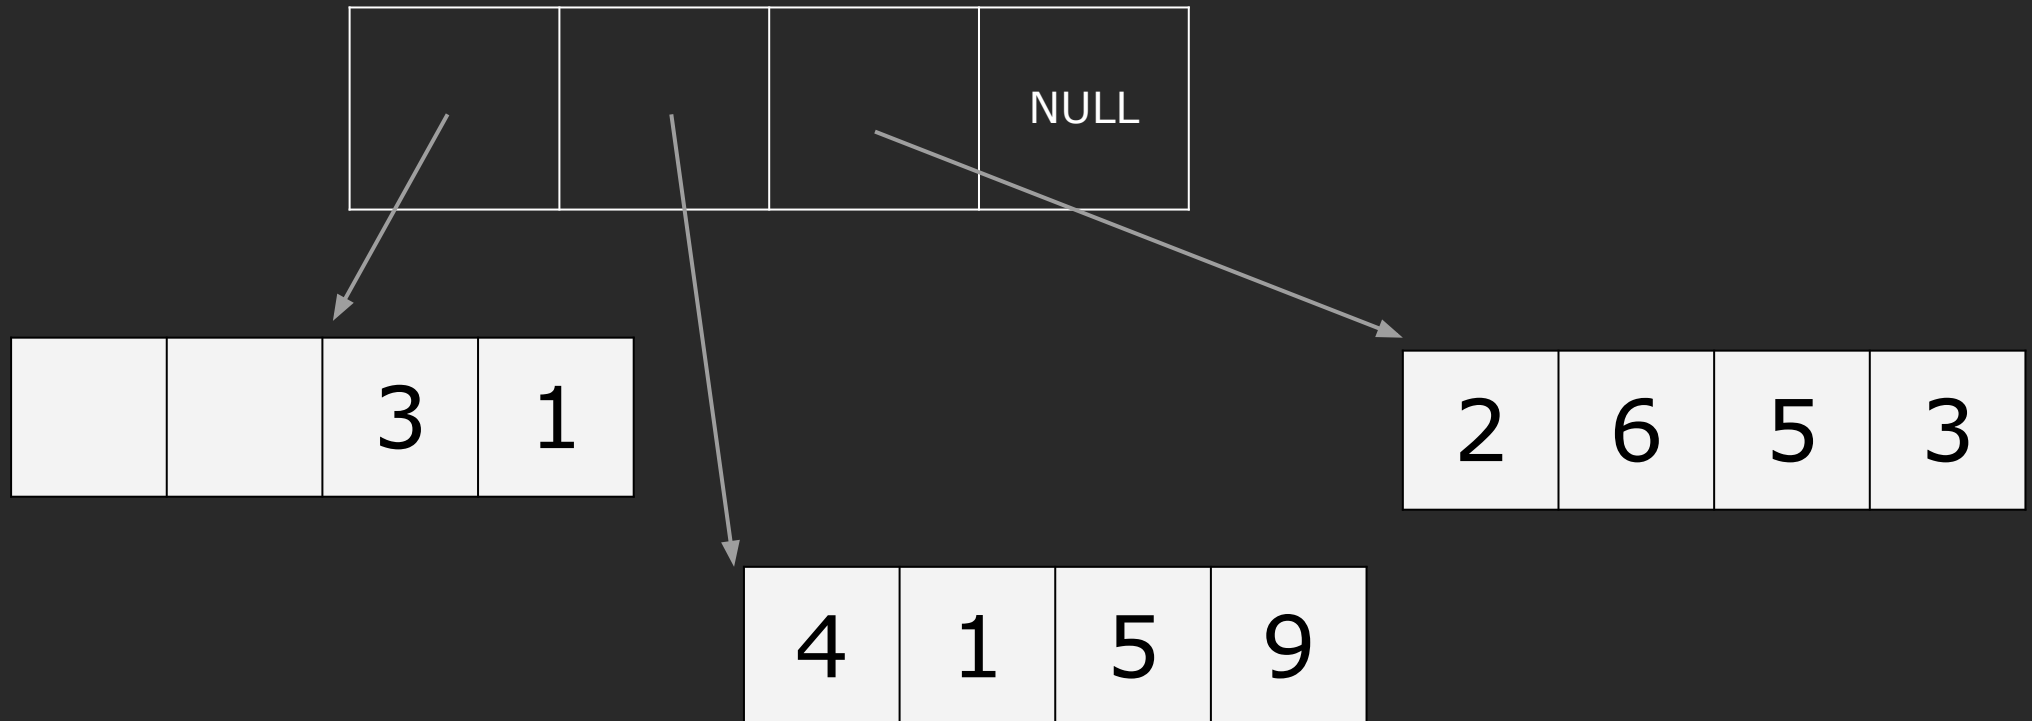
How would you do `push_back(3)`?



Supplemental Material

How does `std::deque<T>` work?

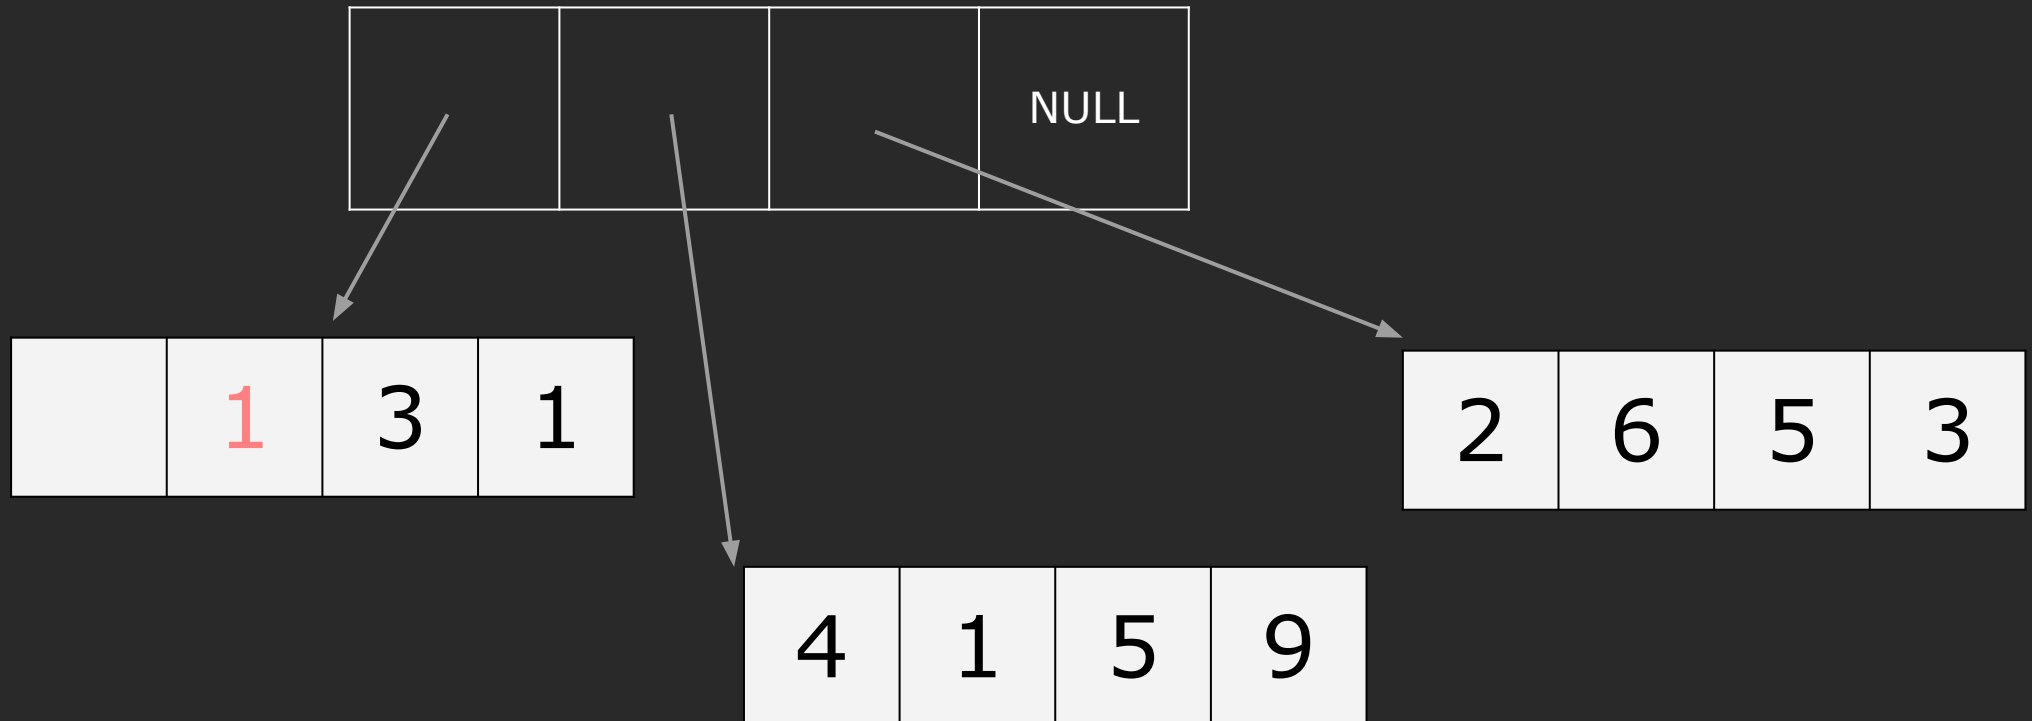
How would you do `push_front(1)`?



Supplemental Material

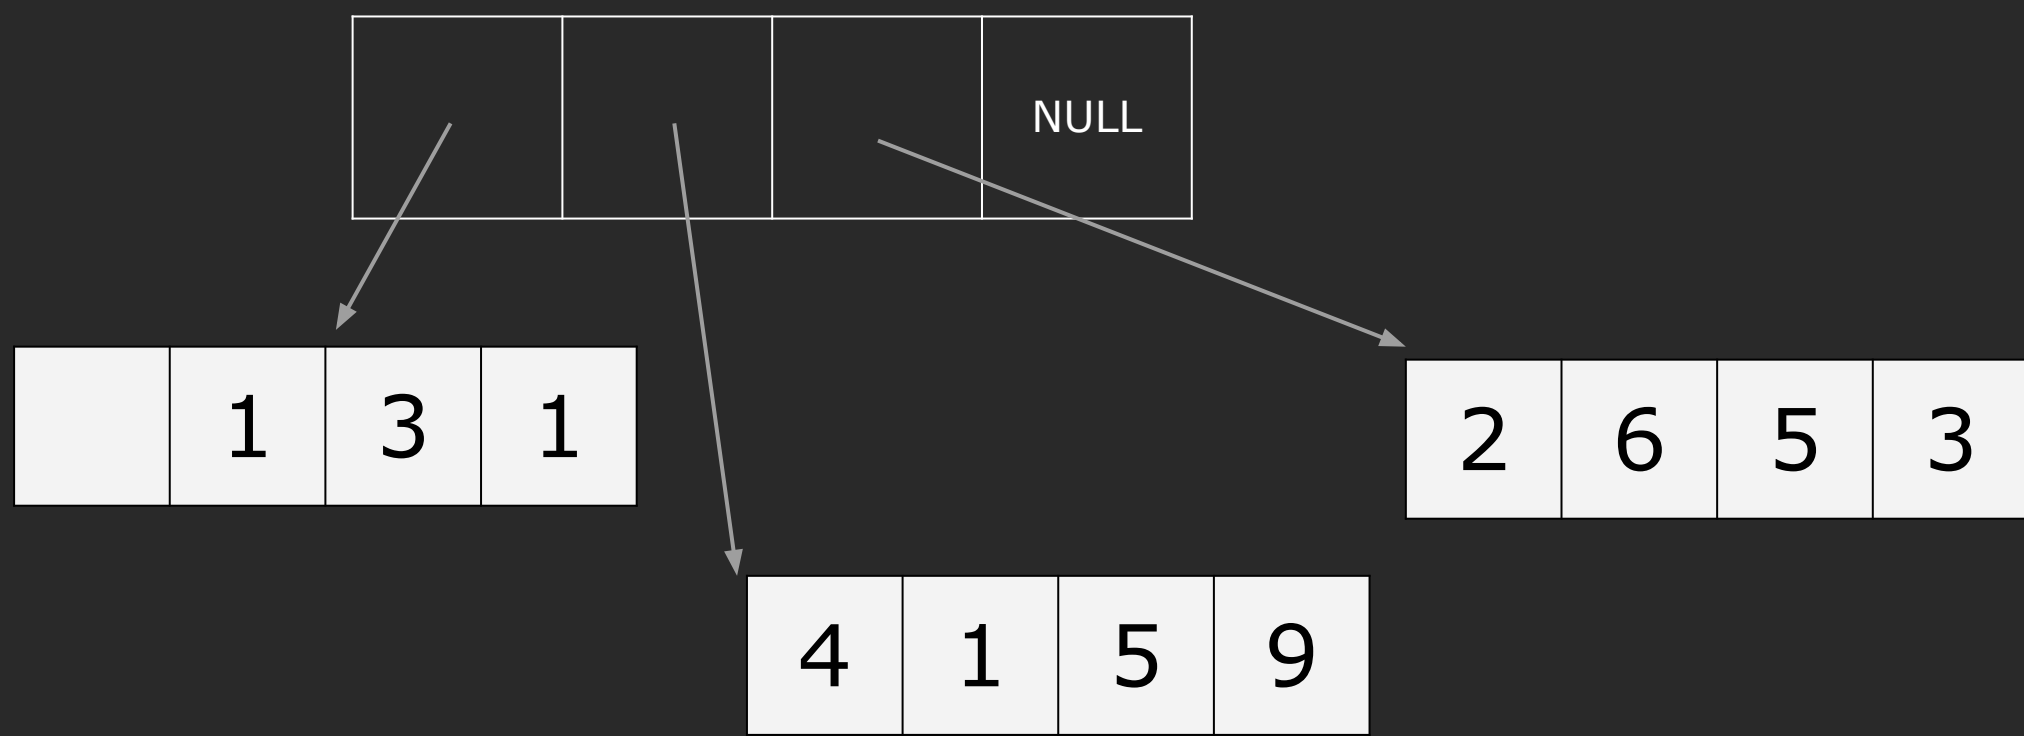
How does `std::deque<T>` work?

How would you do `push_front(1)`?



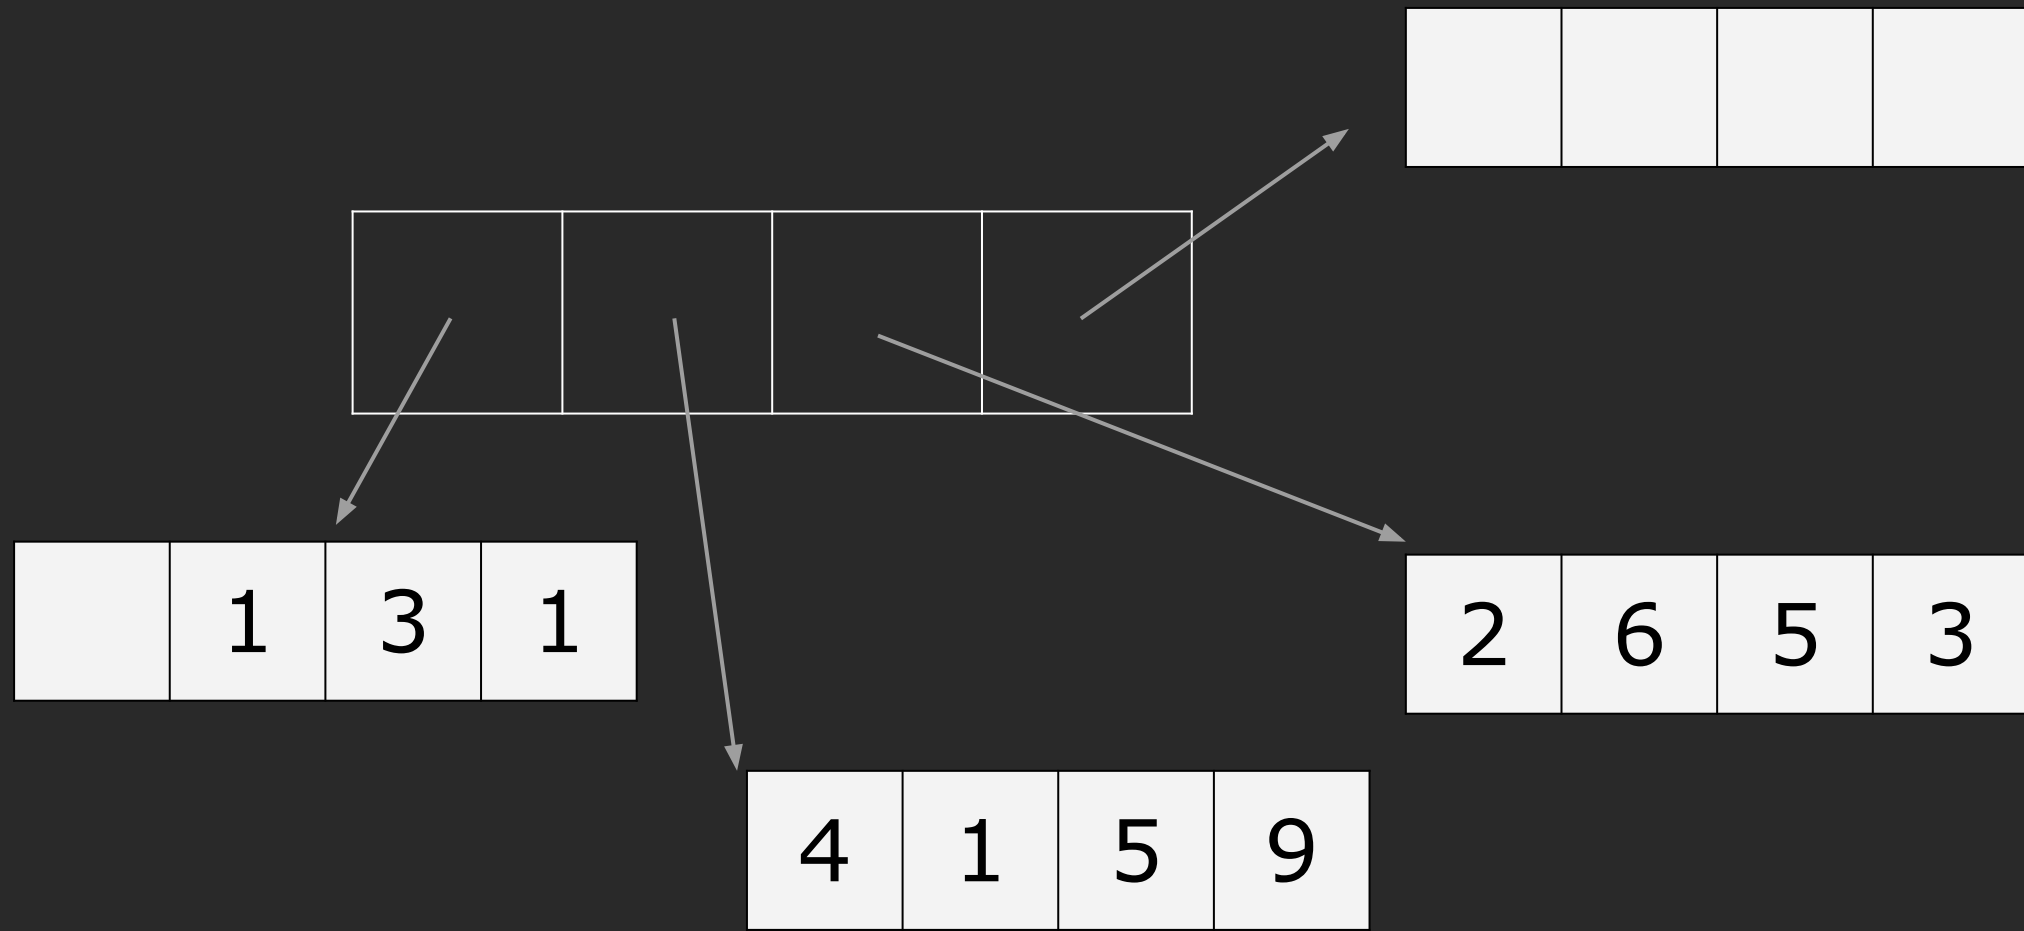
How does `std::deque<T>` work?

How would you do `push_back(7)`?



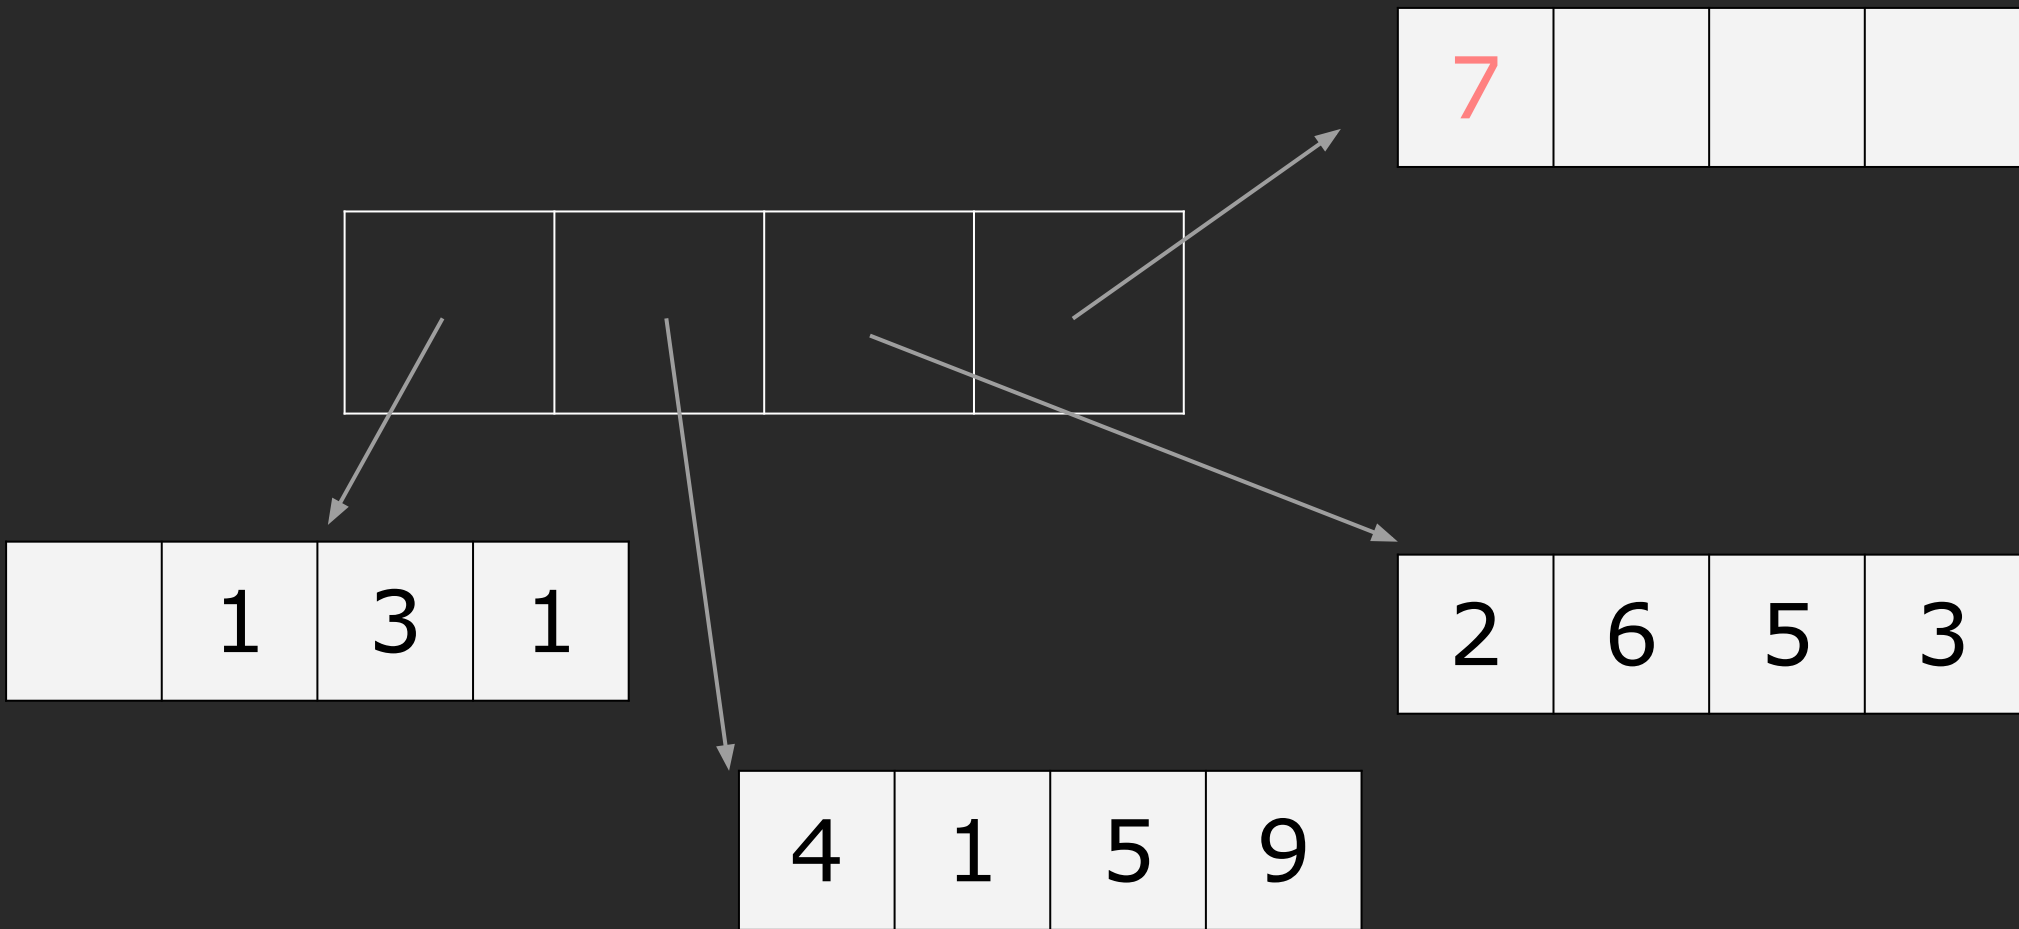
How does `std::deque<T>` work?

How would you do `push_back(7)`?



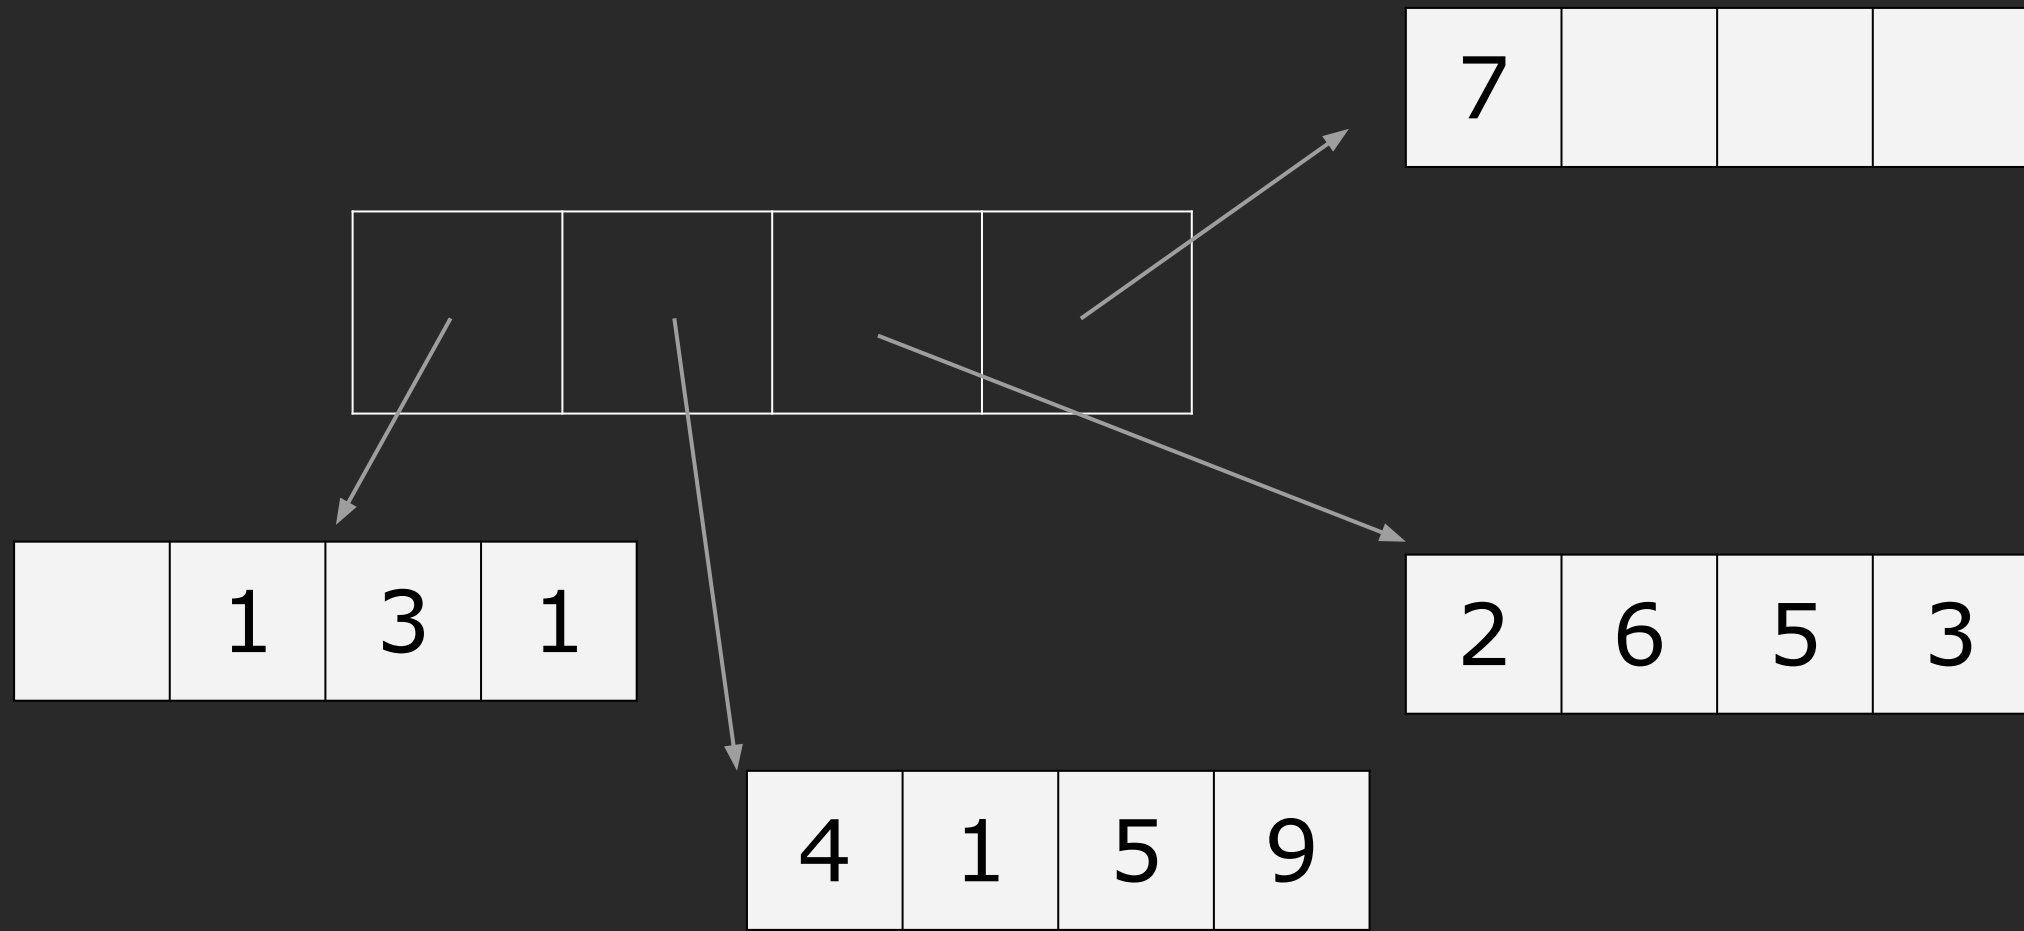
How does `std::deque<T>` work?

How would you do `push_back(7)`?



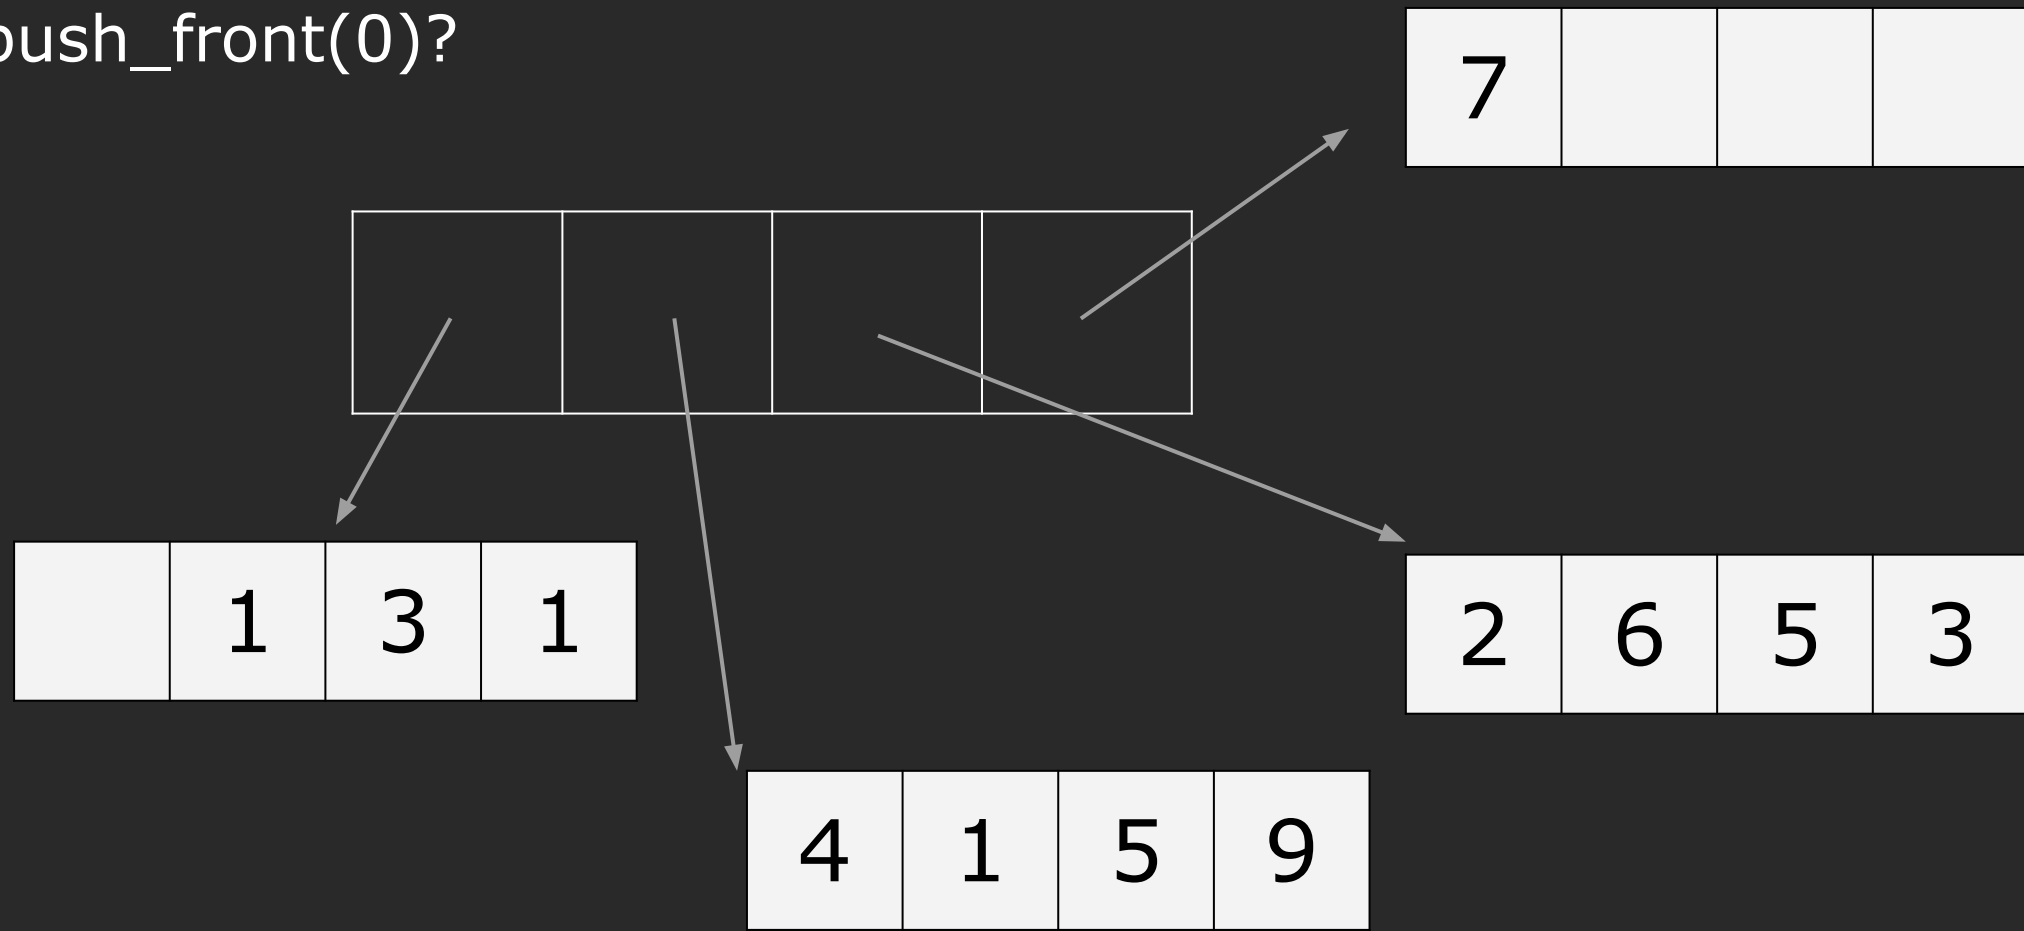
How does `std::deque<T>` work?

How would you do `push_back(7)`?



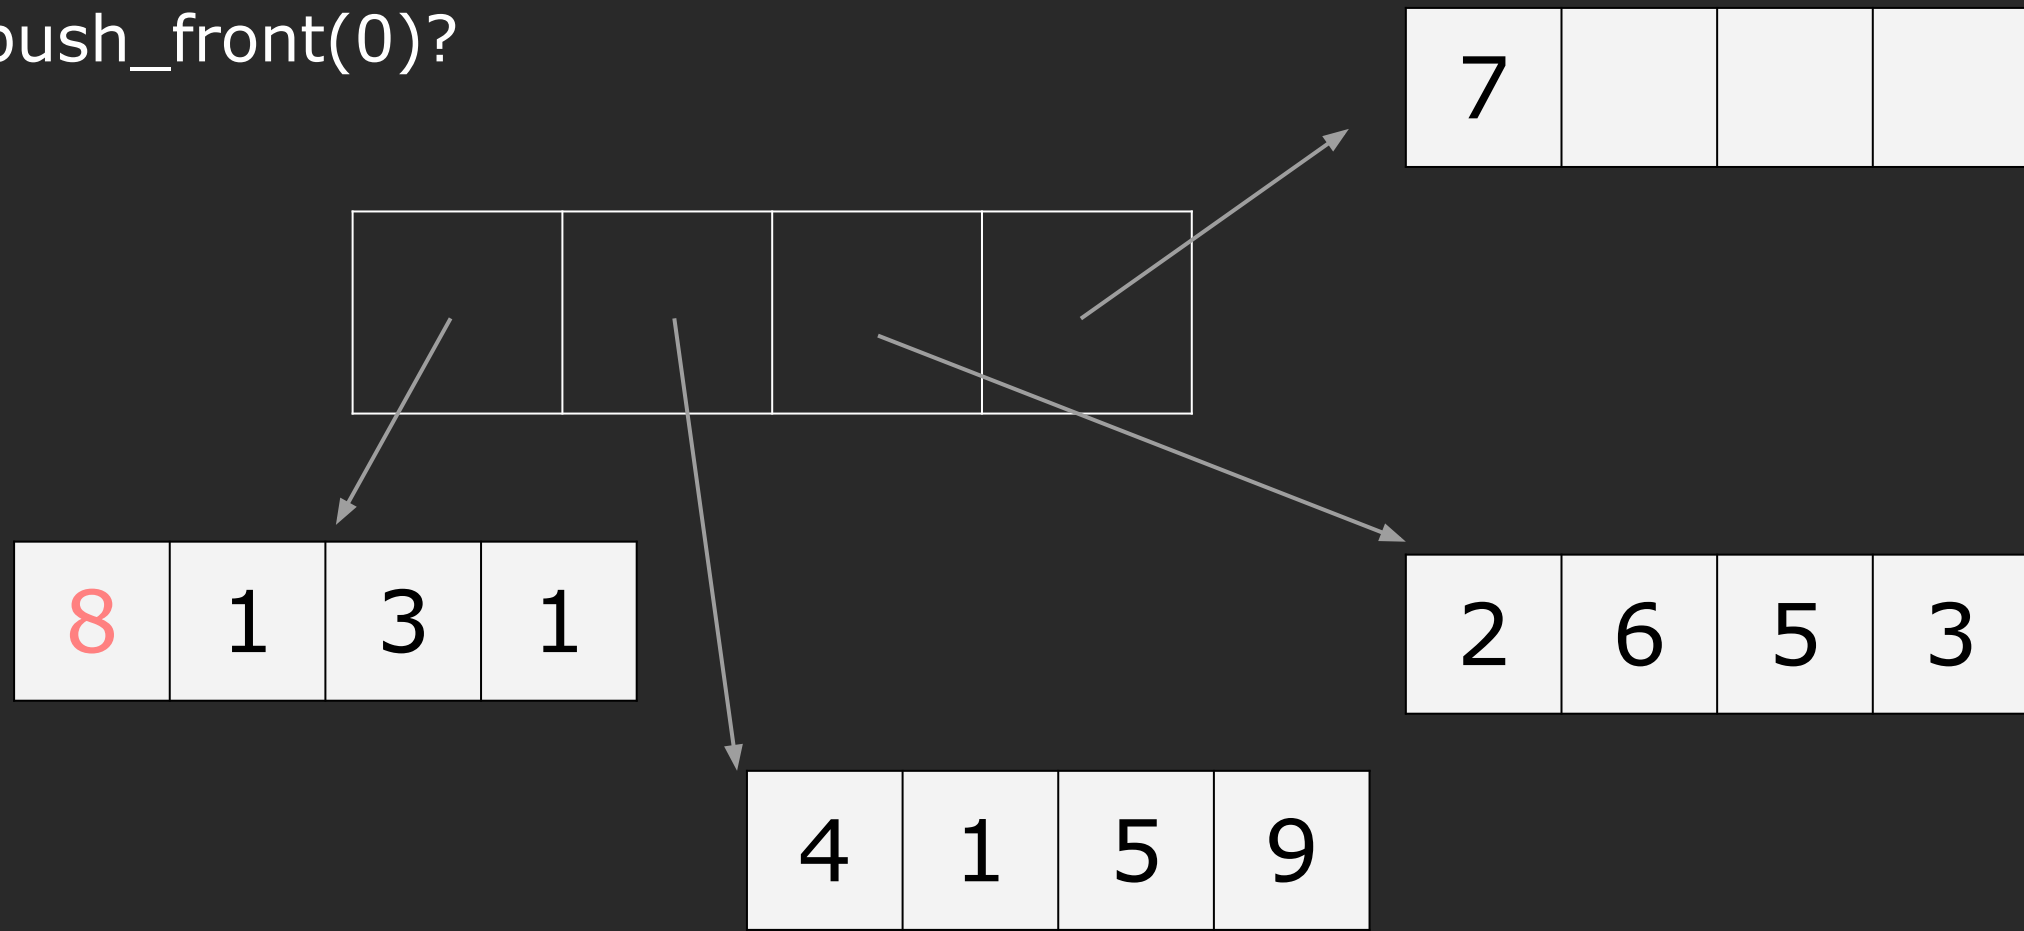
How does `std::deque<T>` work?

How would you do `push_front(8)`
then `push_front(0)`?




How does std::deque<T> work?

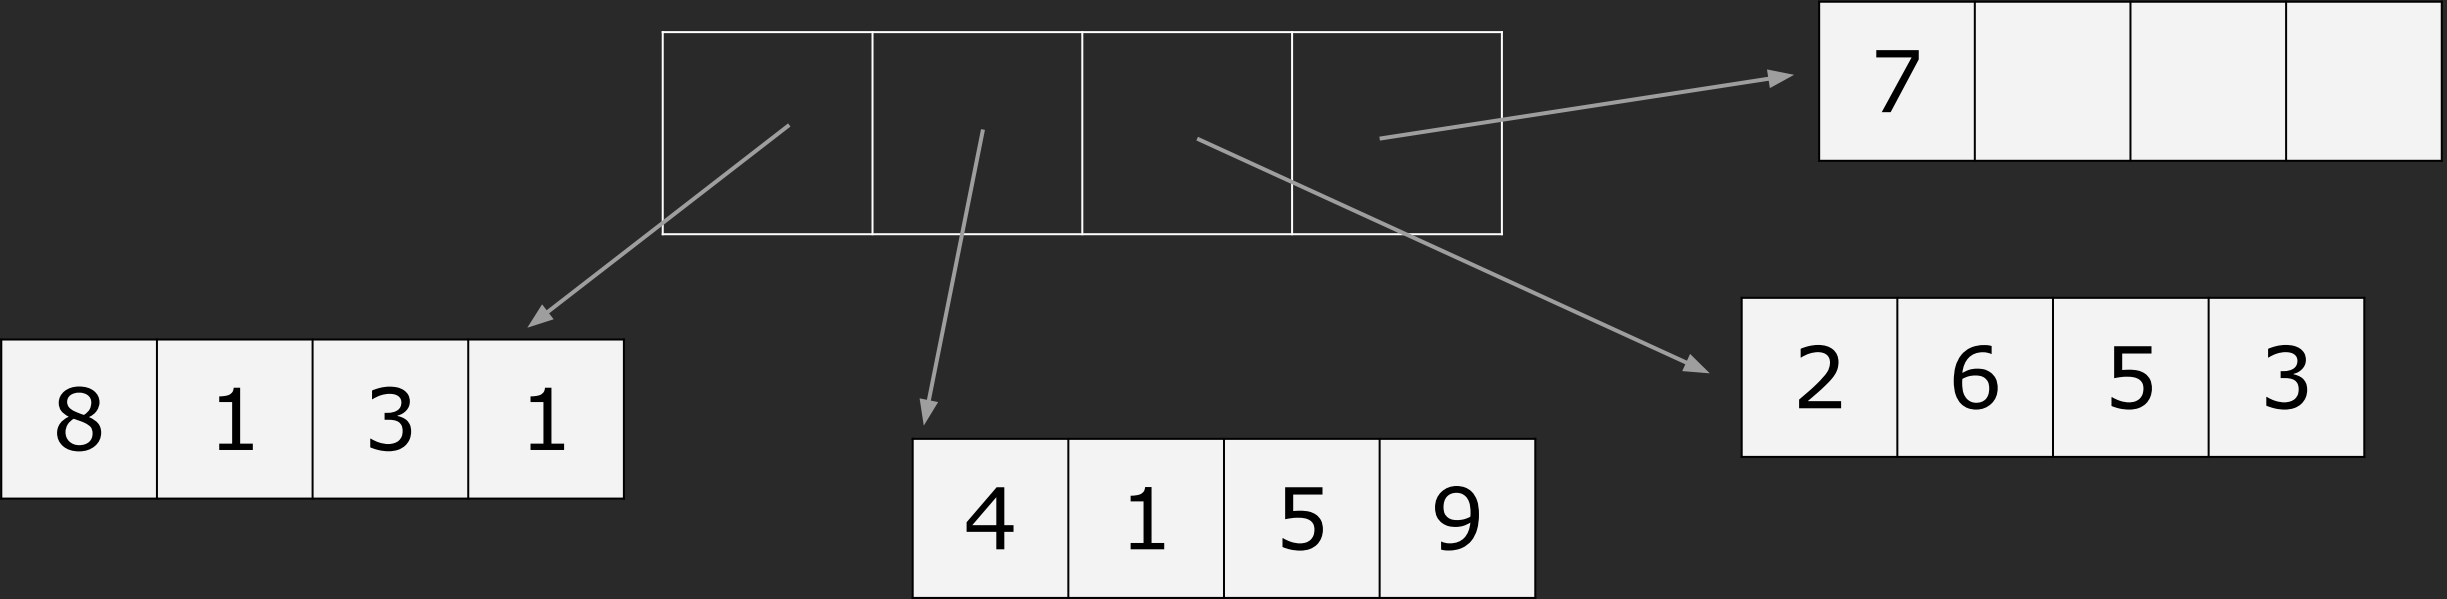
How would you do push_front(8)
then push_front(0)?



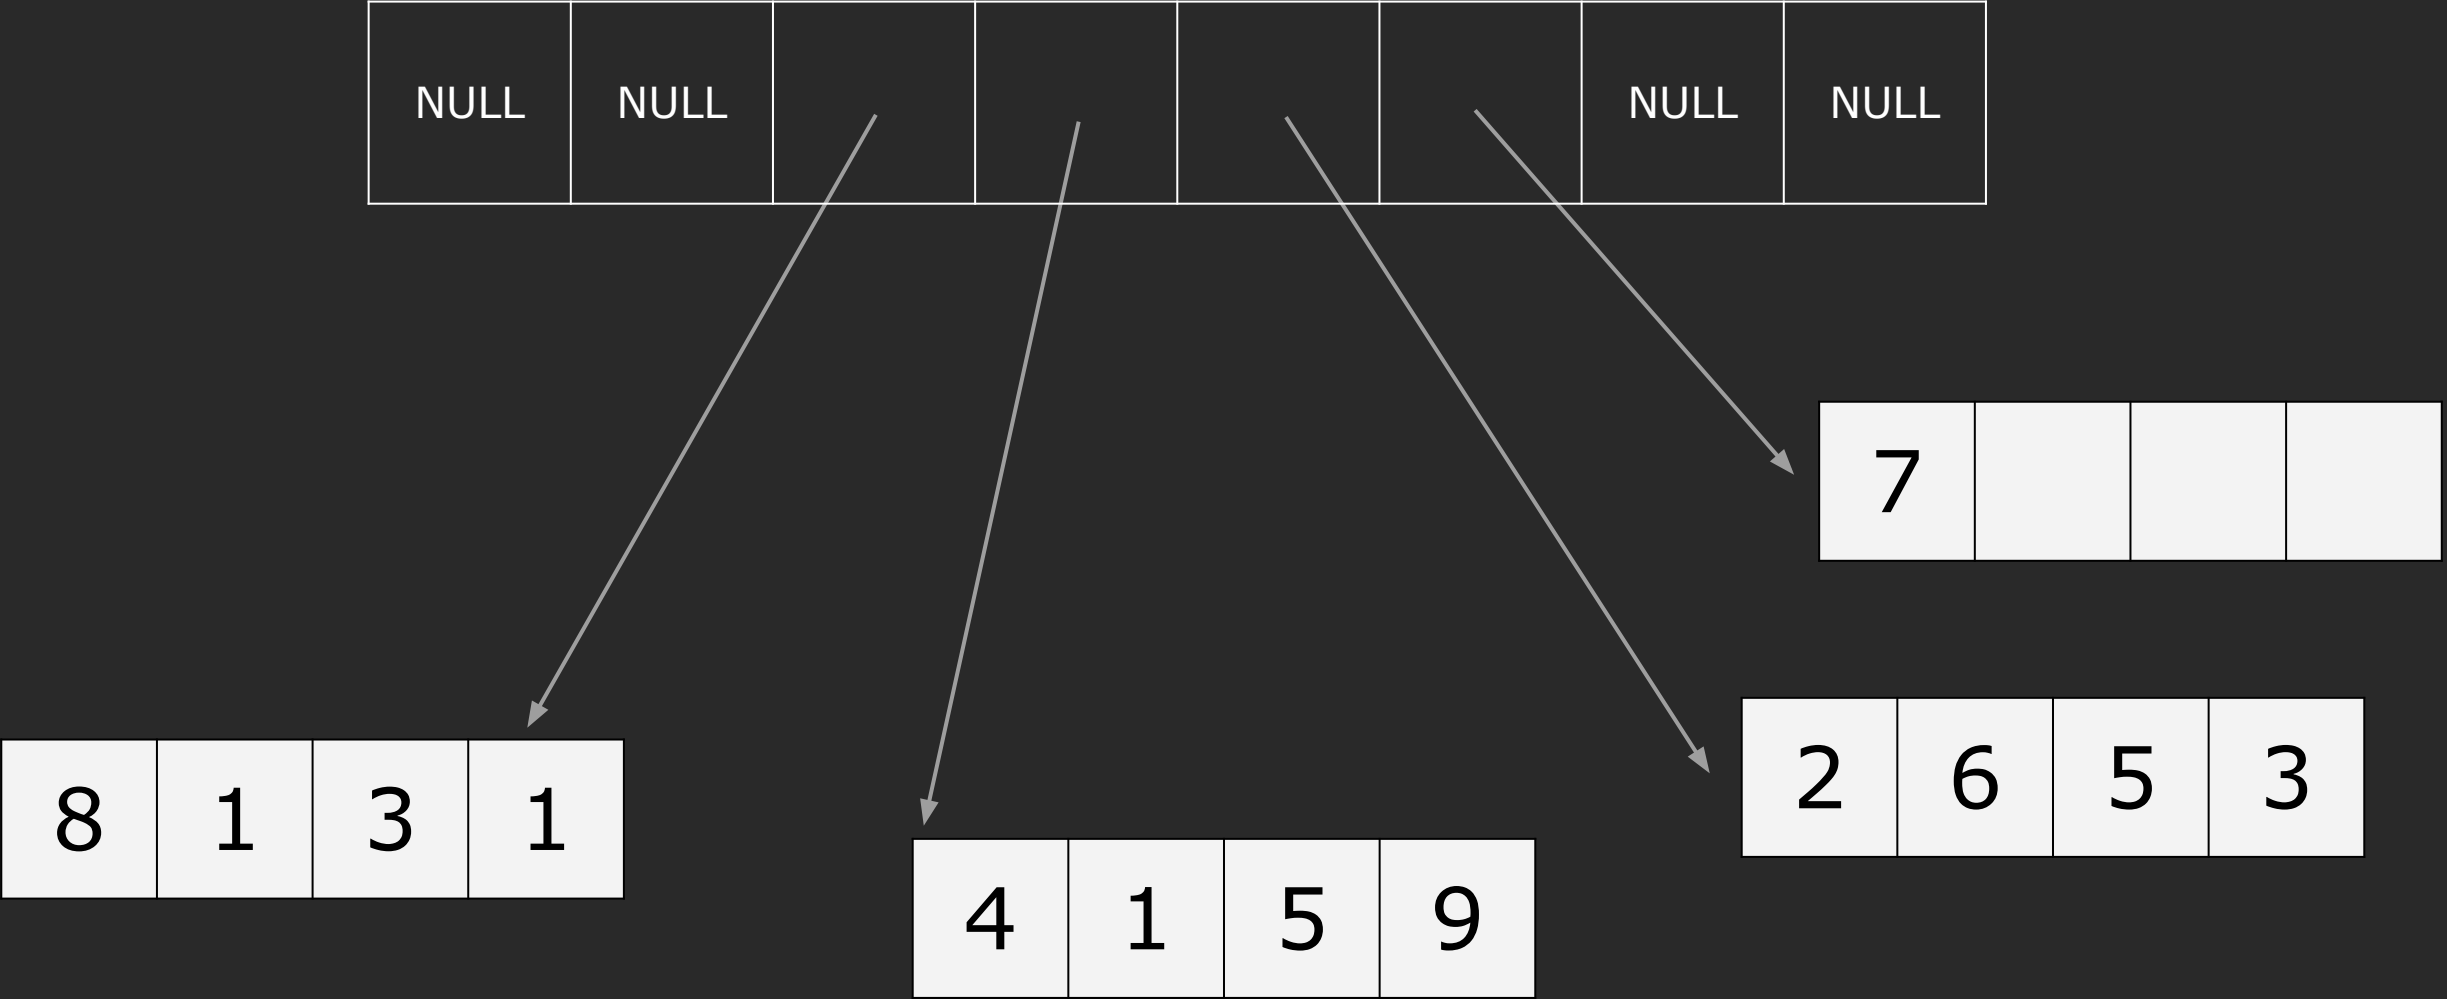
How does `std::deque<T>` work?



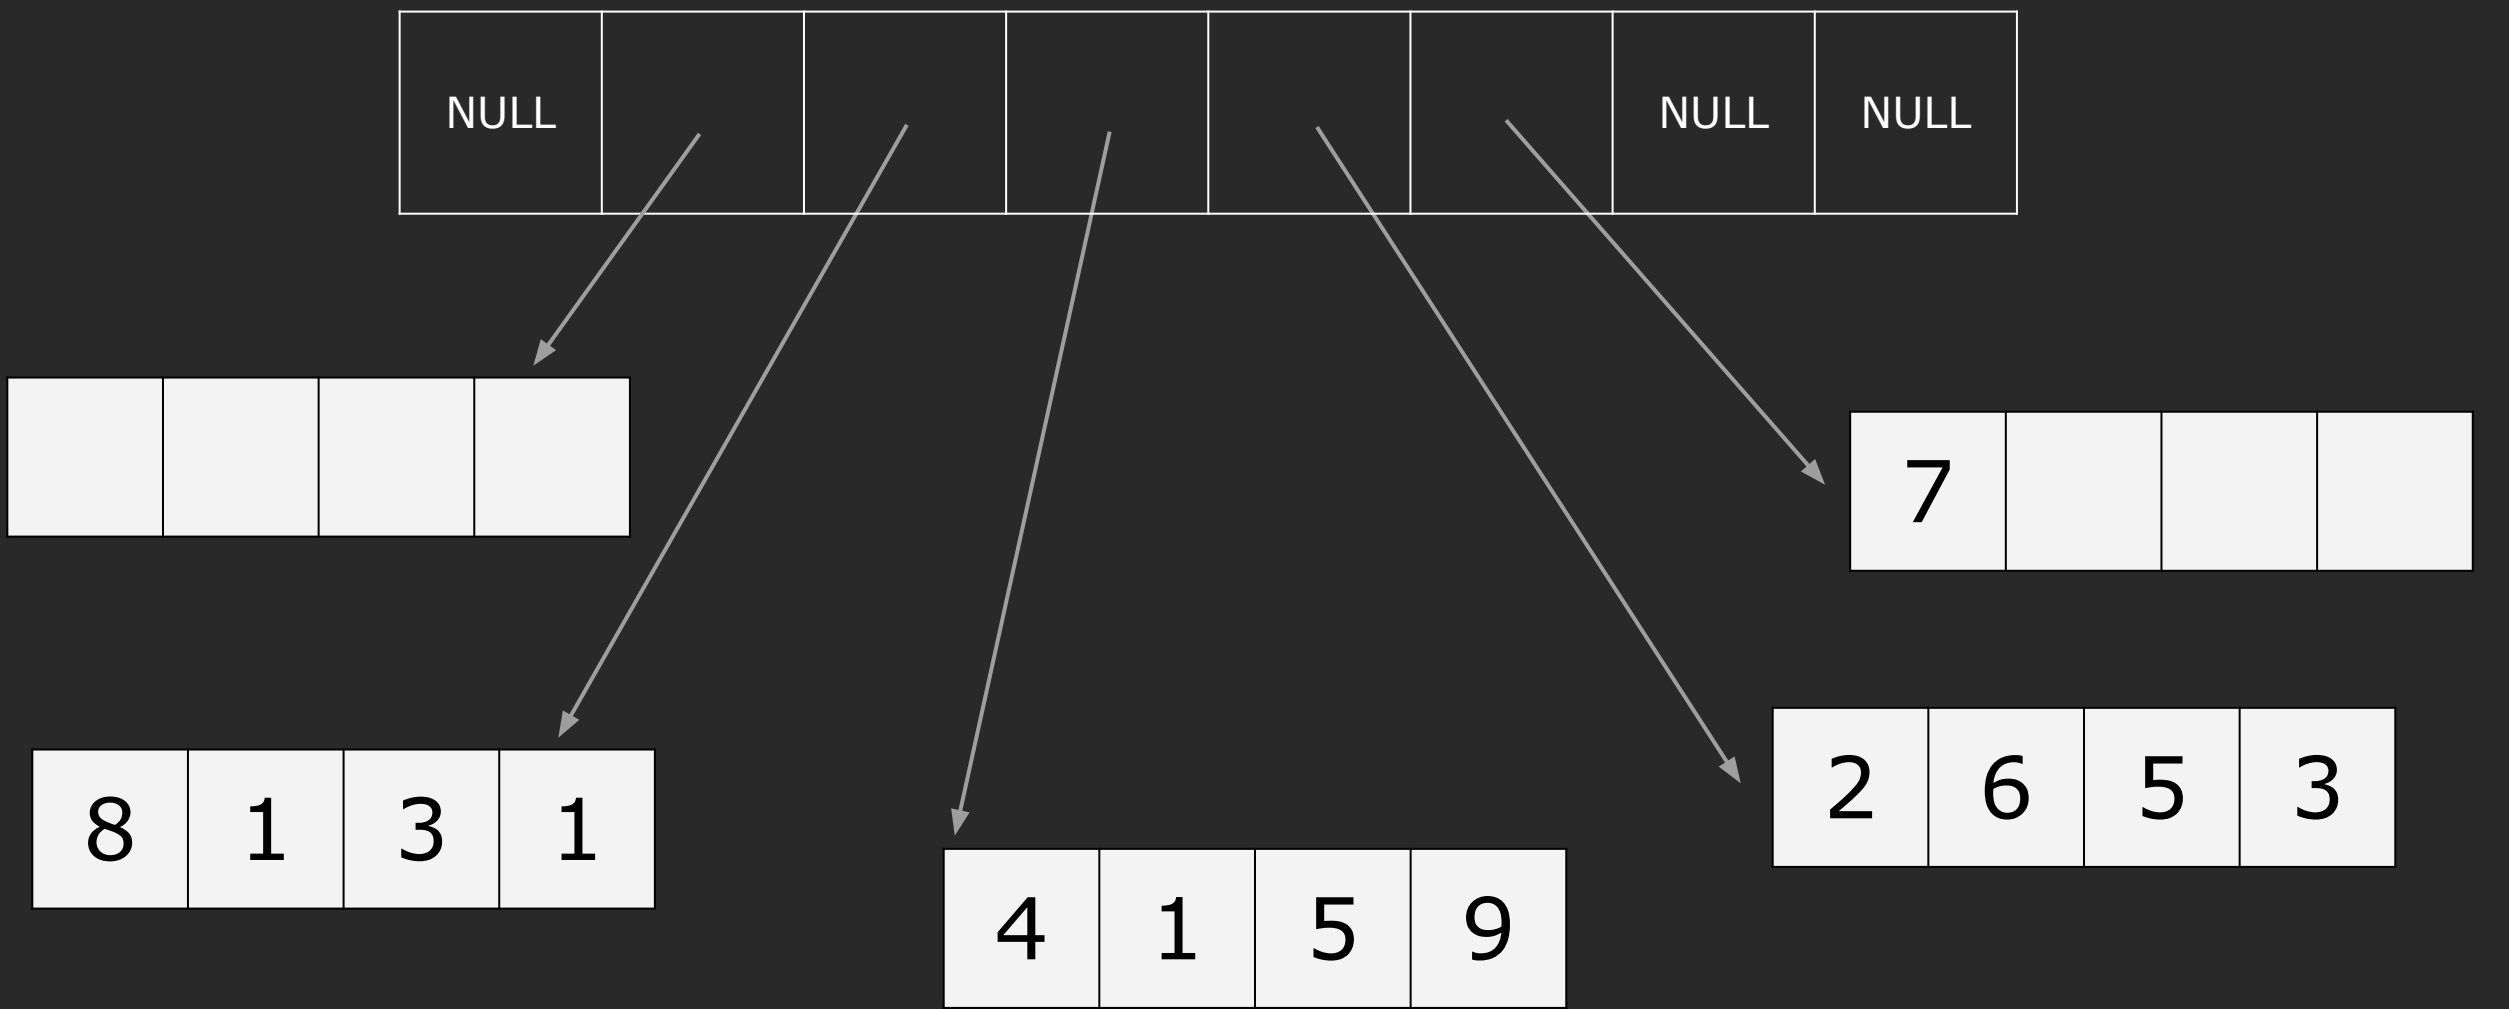
How does `std::deque<T>` work?



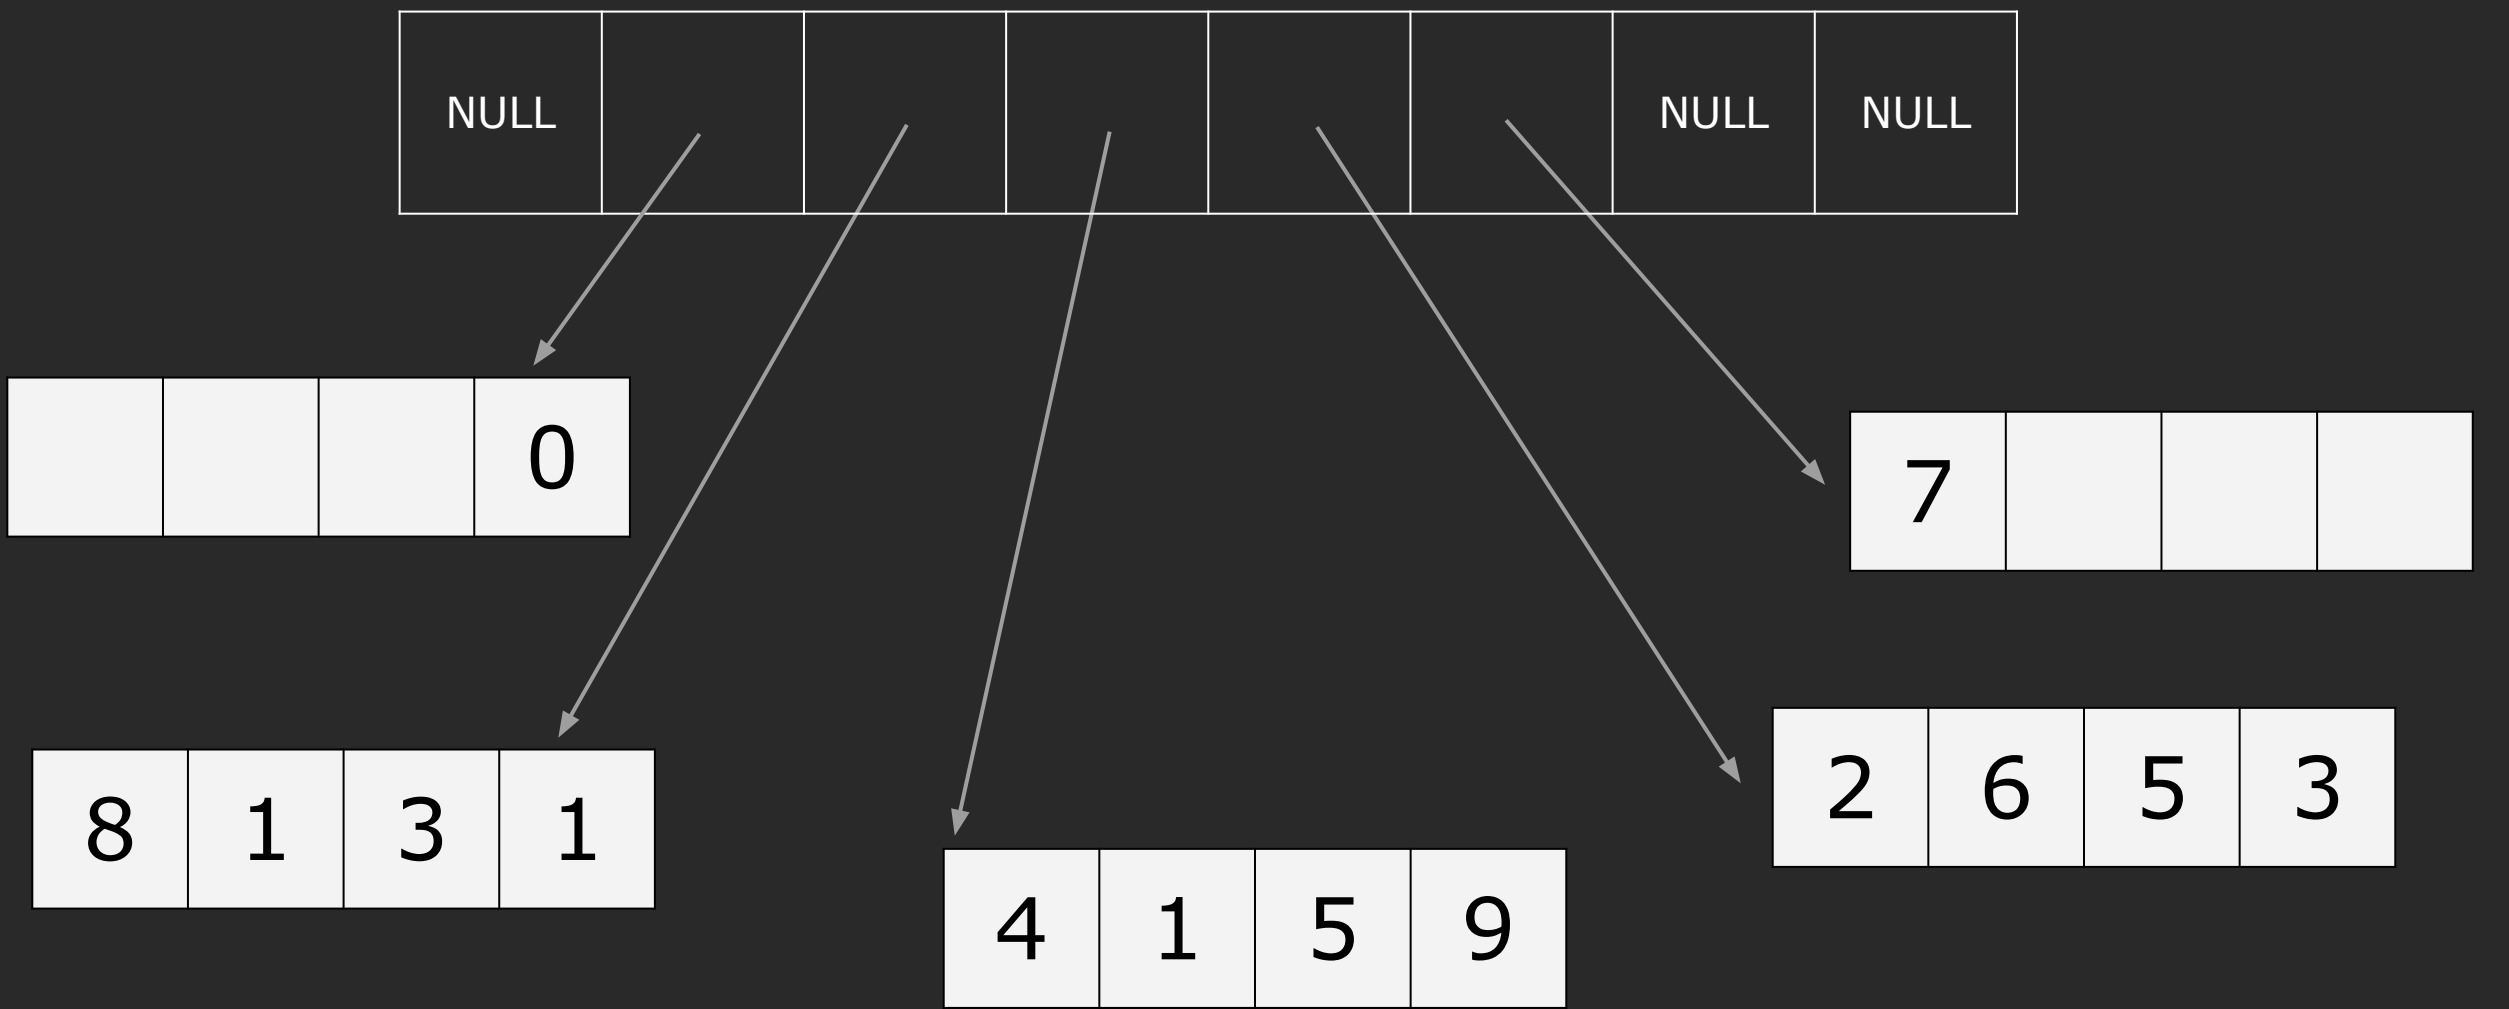
How does `std::deque<T>` work?



How does `std::deque<T>` work?



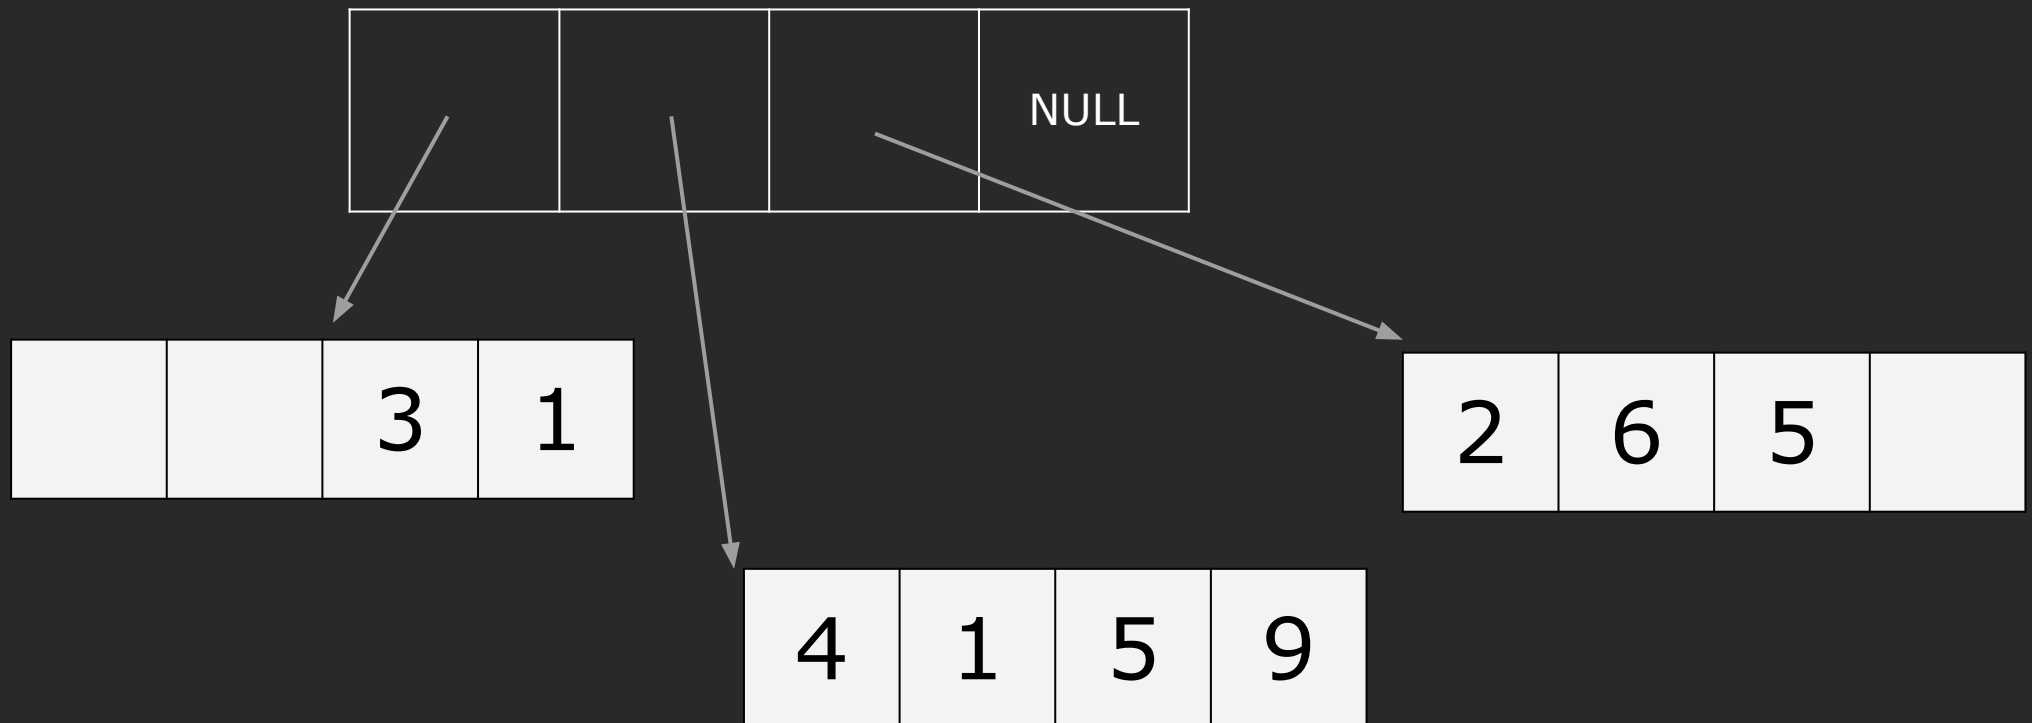
How does `std::deque<T>` work?



How does `std::deque<T>` work?

Difficult thought question that we won't answer in class:

How fast is inserting? Is it better or worse than a `std::vector<T>`?



std::list is kinda like std::stack + std::queue

std::list provides fast removal from the front and end
but you can't access any elements in the middle.

```
1. std::list<int> list{5, 6};    // {5, 6}
2. list.push_front(3);
3. list.pop_back(4);
```

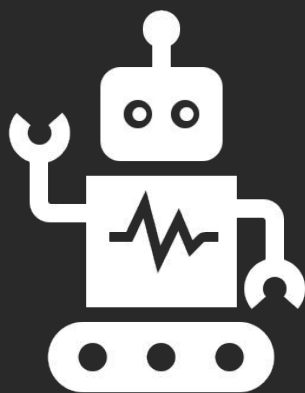
Sidenote: usually a doubly-linked list. There's also a forward_list that's a singly-linked list.

when to use which sequence container?

	std::vector	std::deque	std::list
Indexed Access	Super Fast	Fast	Impossible
Insert/remove front	Slow	Fast	Fast
Insert/remove back	Super Fast	Very Fast	Fast
Ins/rem elsewhere	Slow	Fast	Very Fast
Memory	Low	High	High
Splicing/Joining	Slow	Very Slow	Fast
Stability (Iterators, concurrency)	Poor	Very Poor	Good

Sidenote: color-wise vector might not look great, but remember that indexed access and inserting to the back are the most common uses of sequence containers.

Sidenote: don't take what I say for granted. Run the sample code to test it out yourself!



Example

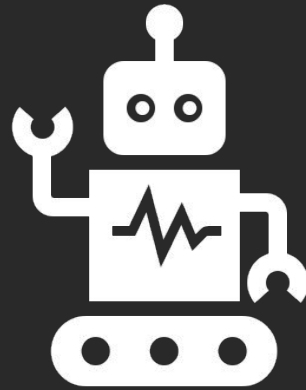
`std::vector` vs. `std::deque` speed

Summary from the ISO Standard

*“vector is the type of sequence that should be used by **default**...
deque is the data structure of choice when most insertions and
deletions take place **at the beginning or at the end** of the
sequence.”*

— C++ ISO Standard (section 23.1.1.2):





Questions

Answer 2 questions.

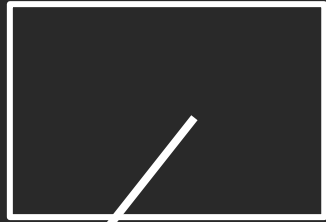
Summary of Sequence Containers

vector: use for most purposes

deque: frequent insert/remove at front

list: very rarely - if need splitting/joining

how is a vector implemented?



internally, a vector consists of an fixed-size array.
the array is automatically resized when necessary.

size = number of elements in the vector

capacity = amount of space saved for the vector

0	3	1	4	1	5	9	2	6		
---	---	---	---	---	---	---	---	---	--	--

best practices: if possible, reserve before insert

What's the best way to create a vector of the first 1,000,000 integers?

```
1. std::vector<int> vec;  
2.  
3. for (size_t i = 0; i < 1000000; ++i) {  
4.     vec.push_back(i);  
5. }
```

Problem: internally the array is resized and copied many times.

best practices: if possible, reserve before insert

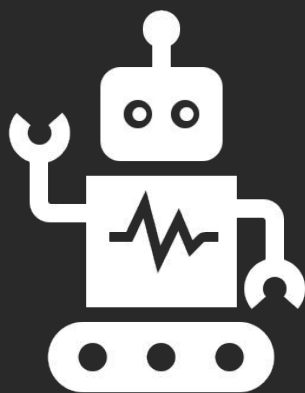
What's the best way to create a vector of the first 1,000,000 integers?

```
1. std::vector<int> vec;  
2. vec.reserve(1000000);  
3. for (size_t i = 0; i < 1000000; ++i) {  
4.     vec.push_back(i);  
5. }
```

Other best practices that we won't go over.

1. Consider using `shrink_to_fit` if you don't need the memory.
2. Call `empty()`, rather than check if `size() == 0`.
3. Don't use `vector<bool>` ("noble failed experiment")
4. A ton of other stuff after we talk about iterators!

If curious, ask us after class!



Example

CS 106B vector examples, using the STL vector/deque/list
Some performance analysis

Summary of Supplemental Material

It's easy to write inefficient code.
Know about the common pitfalls - prevent as much resizing as much as possible.

Container Adaptors

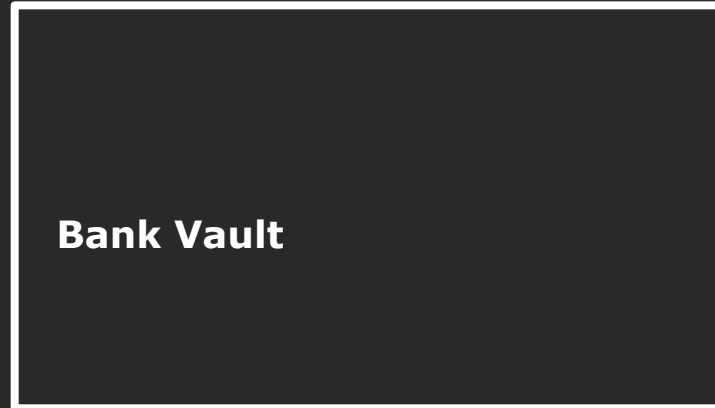
1. What is a container adaptor?
2. `std::stack` and `std::queue`

What is a wrapper (in general)?

A wrapper for an object changes how external users can interact with that object.

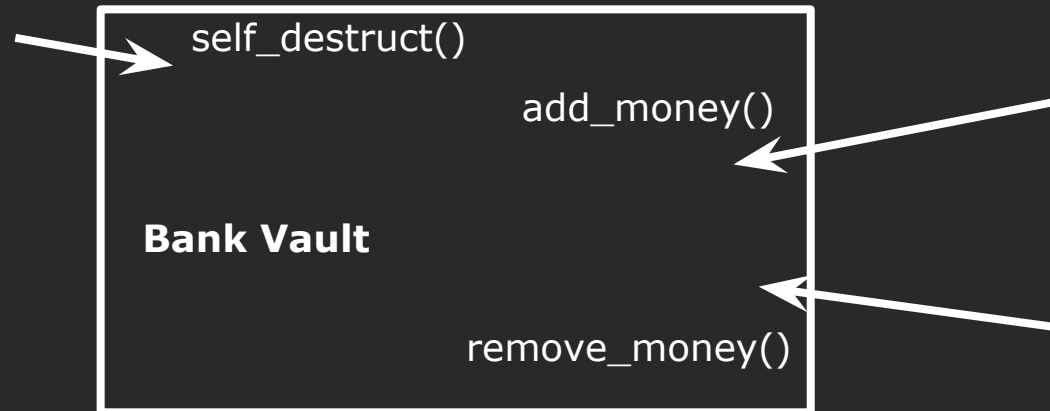
What is a wrapper (in general)?

Here is a bank vault.



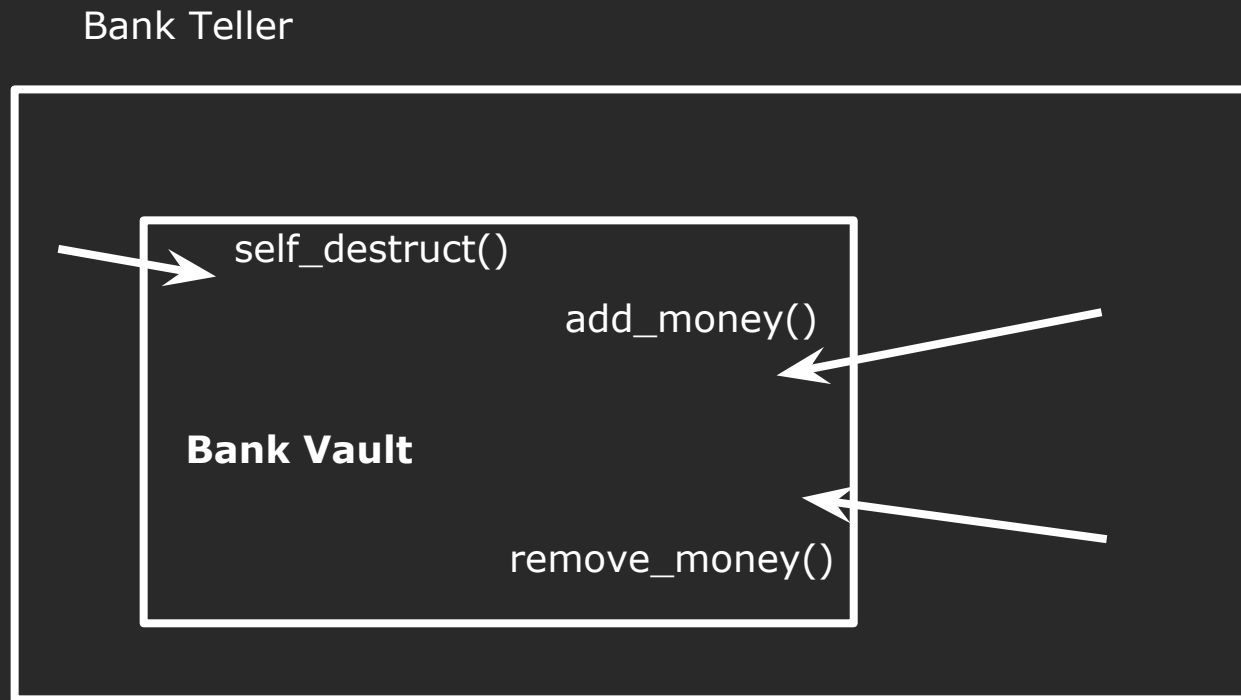
What is a wrapper (in general)?

There are many ways you can interact with a bank vault. It would be bad if people outside the bank could interact in these manners freely.



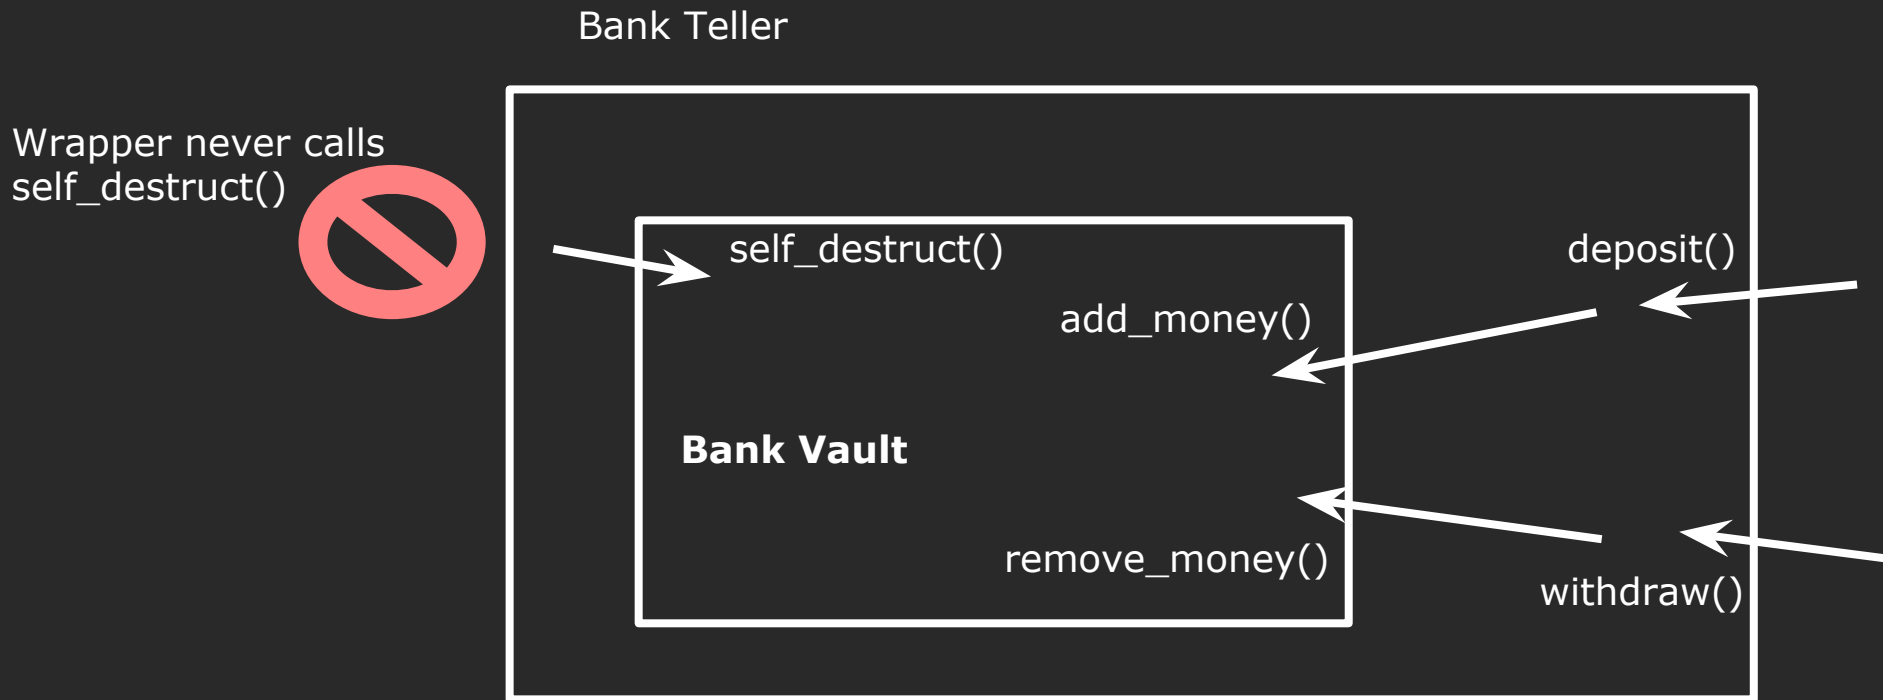
What is a wrapper (in general)?

The bank teller limits your access to the bank value.



What is a wrapper (in general)?

The bank teller is in charge of forwarding your request to the actual bank vault itself.



How do you design a stack?

Container adaptors provide a different interface for sequence containers. You can choose what the underlying container is!

How do you design a stack?

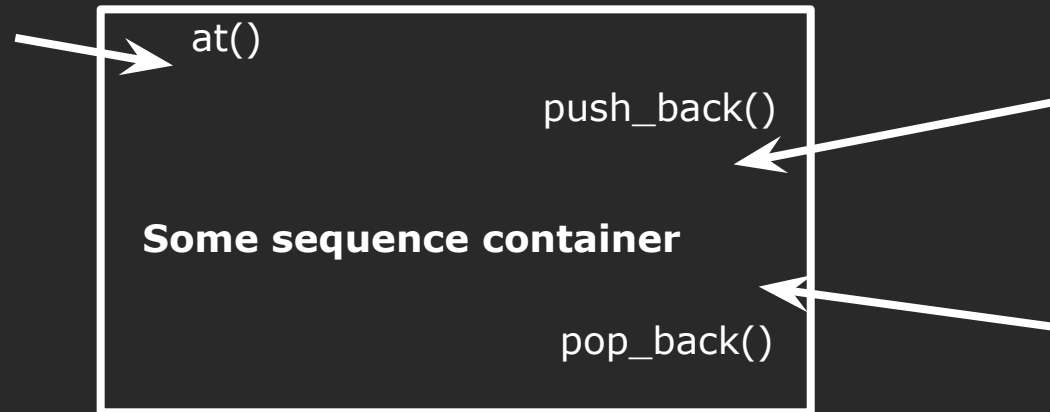
Container adaptors provide a different interface for sequence containers. You can choose what the underlying container is!



Some sequence container

How do you design a stack?

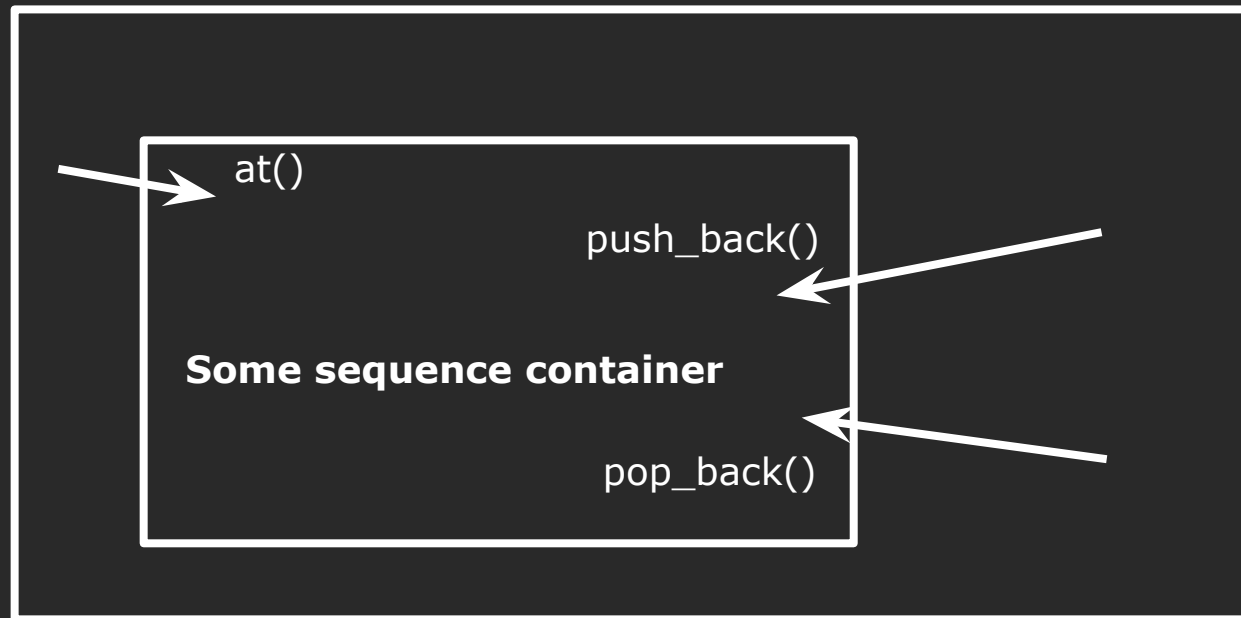
Container adaptors provide a different interface for sequence containers. You can choose what the underlying container is!



How do you design a stack?

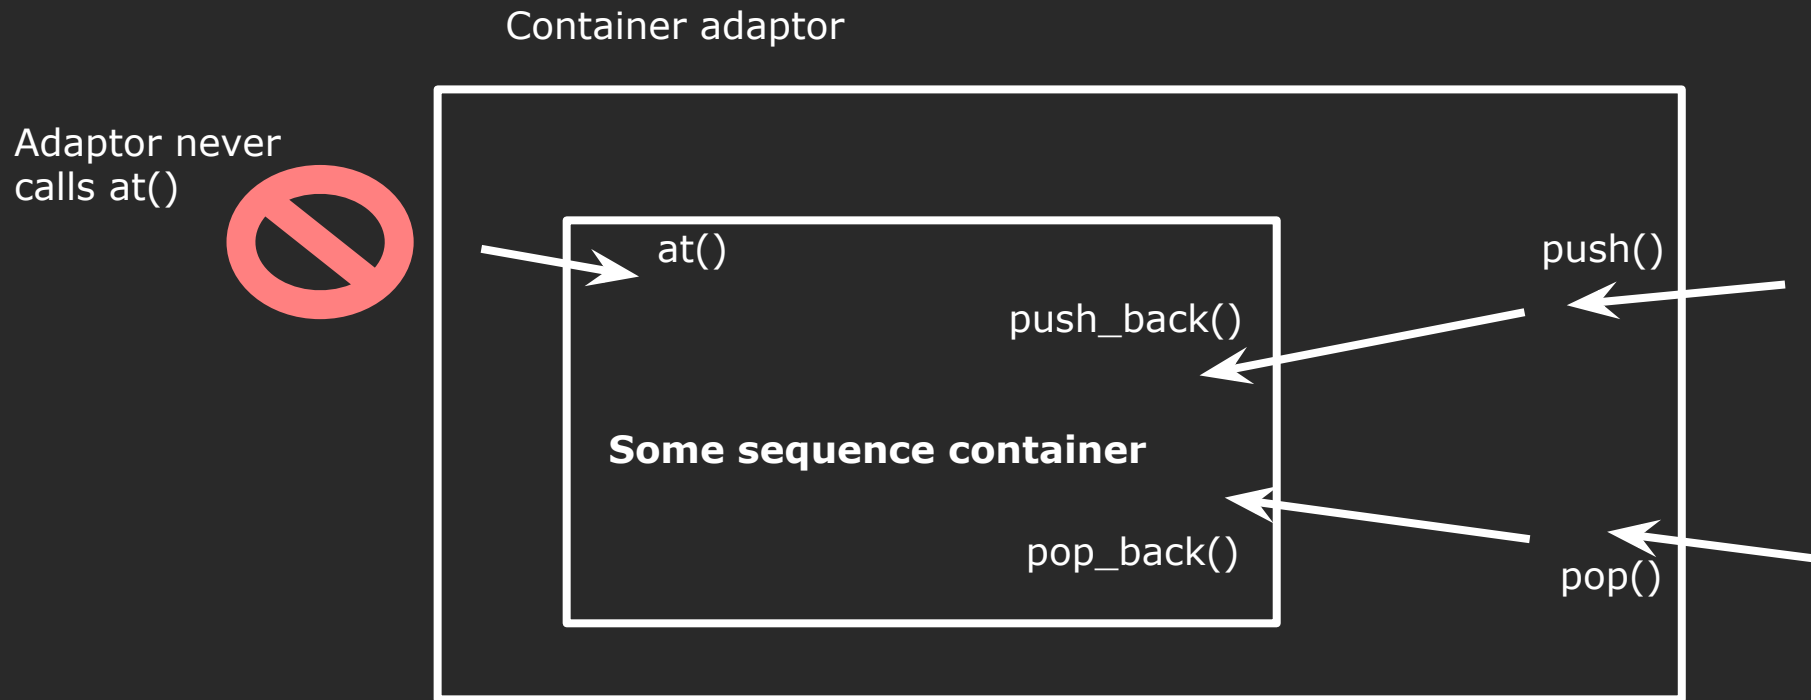
Container adaptors provide a different interface for sequence containers. You can choose what the underlying container is!

Container adaptor



How do you design a stack?

Container adaptors provide a different interface for sequence containers. You can choose what the underlying container is!





Quiz Question

Which data structure should we use to implement a stack? How about a queue?

- A. `std::vector<T>`
- B. `std::deque<T>`
- C. Both are equally good
- D. Both are equally bad



Answer on Poll:

Which data structure should we use to implement a stack? How about a queue?

A. `std::vector<T>`

B. `std::deque<T>`

C. Both are equally good

D. Both are equally bad

Answer on chat:

Why?

Concrete Example

std::stack

Defined in header `<stack>`

```
template<
    class T,
    class Container = std::deque<T>
> class stack;
```

Why deque as opposed to vector or list?

The `std::stack` class is a container adapter that gives the programmer the functionality of a stack - specifically, a LIFO (last-in, first-out) data structure.

The class template acts as a wrapper to the underlying container - only a specific set of functions is provided. The stack pushes and pops the element from the back of the underlying container, known as the top of the stack.

std::queue

Defined in header `<queue>`

```
template<
    class T,
    class Container = std::deque<T>
> class queue;
```

No surprise

The `std::queue` class is a container adapter that gives the programmer the functionality of a queue - specifically, a FIFO (first-in, first-out) data structure.

The class template acts as a wrapper to the underlying container - only a specific set of functions is provided. The queue pushes the elements on the back of the underlying container and pops them from the front.

Concrete Example

std::stack

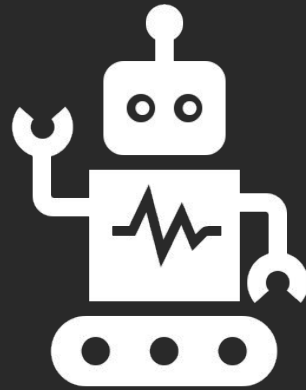
Defined in header `<stack>`

```
template<
    class T,
    class Container = std::deque<T>
> class stack;
```

The `std::stack` class is a container adapter that gives the programmer the functionality of a stack - specifically, a LIFO (last-in, first-out) data structure.

The class template acts as a wrapper to the underlying container - only a specific set of functions is provided. The stack pushes and pops the element from the back of the underlying container, known as the top of the stack.

```
std::stack<int> stack_d;           // Container = deque
std::stack<int, std::vector<int>> stack_v; // Container = vector
std::stack<int, std::list<int>> stack_l;  // Container = list
```



Questions

Answer 2 questions.

You'll be using `std::priority_queue` for A1!

`std::priority_queue`

Defined in header `<queue>`

```
template<
    class T,
    class Container = std::vector<T>,
    class Compare = std::less<typename Container::value_type>
> class priority_queue;
```

A priority queue is a container adaptor that provides constant time lookup of the largest (by default) element, at the expense of logarithmic insertion and extraction.

A user-provided Compare can be supplied to change the ordering, e.g. using `std::greater<T>` would cause the smallest element to appear as the `top()`.

Working with a `priority_queue` is similar to managing a `heap` in some random access container, with the benefit of not being able to accidentally invalidate the heap.

CS 106B A5 is basically to write this container adaptor.

Associative Containers

1. `std::set` functions
2. `std::map` functions and auto-insertion
3. type requirements

Stanford vs. STL set: a summary

What you want to do	Stanford Set<int>	std::set<int>
Create an empty set	Set<int> s;	set<int> s;
Add k to the set	s.add(k);	s.insert(k);
Remove k from the set	s.remove(k);	s.erase(k);
Check if k is in the set (* C++20)	if (s.contains(k)) ...	if (s.count(k)) ... if (s.contains(k)) (*)
Check if the set is empty	if (s.isEmpty()) ...	if (s.empty()) ...

Answer on chat: what is different?

Stanford vs. STL set: a summary

What you want to do	Stanford Set<int>	std::set<int>
Create an empty set	<code>Set<int> s;</code>	<code>set<int> s;</code>
Add k to the set	<code>s.add(k);</code>	<code>s.insert(k);</code>
Remove k from the set	<code>s.remove(k);</code>	<code>s.erase(k);</code>
Check if k is in the set (* C++20)	<code>if (s.contains(k)) ...</code>	<code>if (s.count(k)) ...</code> <code>if (s.contains(k)) (*)</code>
Check if the set is empty	<code>if (s.isEmpty()) ...</code>	<code>if (s.empty()) ...</code>

There are functions for size, ==, !=, clear, etc.

STL does not have member functions for subset, difference, union, intersection, etc,
but there are STL algorithms!

Stanford vs. STL map: a summary

What you want to do	Stanford Map <int, char>	std:: map <int, char>
Create an empty map	Map <int, char> m;	map <int, char> m;
Add key k with value v into the map	m.put(k, v);	m.insert({k, v});
Remove key k from the map	m.remove(k);	m.erase(k);
Check if k is in the map (* C++20)	if (m.containsKey(k)) ...	if (m.count(k)) ... if (m.contains(k)) (*)
Check if the map is empty	if (m.isEmpty()) ...	if (m.empty()) ...
Retrieve or overwrite value associated with key k (error if does not exist)	Impossible (put does auto-insert)	char c = m.at(k); m.at(k) = v;
Retrieve or overwrite value associated with key k (auto-insert if DNE)	char c = m[k]; m[k] = v;	char c = m[k]; m[k] = v;

Stanford vs. STL map: a summary

What you want to do	Stanford Map <int, char>	std:: map <int, char>
Create an empty map	<code>Map<int, char> m;</code>	<code>map<int, char> m;</code>
Add key k with value v into the map	<code>m.put(k, v);</code>	<code>m.insert({k, v});</code>
Remove key k from the map	<code>m.remove(k);</code>	<code>m.erase(k);</code>
Check if k is in the map (* C++20)	<code>if (m.containsKey(k)) ...</code>	<code>if (m.count(k)) ...</code> <code>if (m.contains(k)) (*)</code>
Check if the map is empty	<code>if (m.isEmpty()) ...</code>	<code>if (m.empty()) ...</code>
Retrieve or overwrite value associated with key k (error if does not exist)	Impossible (put does auto-insert)	<code>char c = m.at(k);</code> <code>m.at(k) = v;</code>
Retrieve or overwrite value associated with key k (auto-insert if DNE)	<code>char c = m[k];</code> <code>m[k] = v;</code>	<code>char c = m[k];</code> <code>m[k] = v;</code>

STL maps stores `std::pairs`.

The underlying type stored in a

`std::map<K, V>`

is a

`std::pair<const K, V>.`

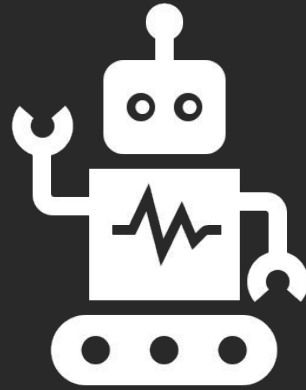
sidenote: why is the key `const`? You should think about this, and you'll need a good answer when A2 rolls around.

Stanford/STL sets + maps require comparison operator

By default, the type (for sets) or key-type (for maps) must have a comparison (<) operator defined for them.

(There is an alternative that we will cover in ~2 lectures)

```
1.  std::set<std::pair<int, int>> set1;  // OKAY - comparable
2.  std::set<std::ifstream> set2;      // ERROR - not comparable
3.
4.  std::map<std::set<int>, int> map1;   // OKAY - comparable
5.  std::map<std::function, int> map2;  // ERROR - not comparable
```



Questions

Answer 4 questions.

Iterating over the elements of a STL set/map.

- Exactly the same as in CS 106B - no modifying the container in the loop!
- The elements are ordered based on the operator< for element/key.
- Because maps store pairs, each element m is an std::pair that you can use structured binding on.

```
1. for (const auto& element : s) {  
2.     // do stuff with key  
3. }  
4.  
5. for (const auto& [key, value] : m) {  
6.     // do stuff with key, value  
7. }
```

unordered_map and unordered_set

Each STL `set/map` comes with an `unordered` sibling. Their usage is the same, with two differences:

- Instead of a comparison operator, the element (set) or key (map) must have a `hash function` defined for it (you'll learn more in A2).
- The `unordered_map/unordered_set` is generally faster than `map/set`.

Don't worry too much about these just yet - you'll implement `unordered_map` in assignment 2!

multimap, multiset (+ unordered siblings)

Each STL set/map (+ unordered_set/map) comes with an multi- cousin. Their usage is the same, with two differences:

- You can have multiple of the same element (set) or key (map).
- insert, erase, and retrieving behave differently (since they may potentially have to retrieve multiple elements/values).

I've actually never used these before. Let us know if you actually ever use one. Otherwise, let's not spend too much time on this.

Summary of Associative Containers

set: membership

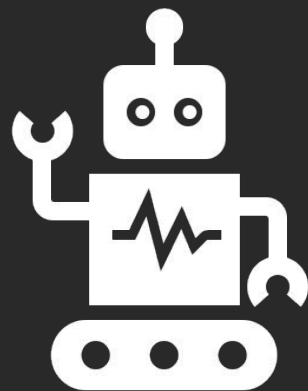
map: key/value pairs

set/map require comparison operator

map stores `std::pairs`

auto-insertion = useful but be careful

unordered/multi stuff - know they exist



Deep Dive into Documentation

map and set

We've run into a few problems with STL containers

Concrete things we haven't been able to do yet...

- how to insert an element at a certain position in a vector?
- how do we create sublists of a vector? how about for a set?
- how do we iterate over all the elements in an associative container?

Ideas we want to move toward...


- can I write a function that works for any container?
- can I operate over a container while completely ignoring the container?

Why does CS 106B not teach the Standard Libraries?

(it's not a fluke. there's a genius pedagogical reason behind it)

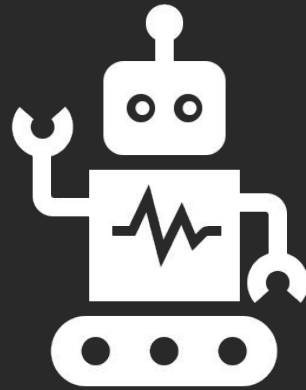
Today's Lecture in Bullet Points

- `std::vector` methods are very similar to Stanford Vector
- `std::vector` indexing has `[]` and `at()` functions. only `at()` does error-checking
- `std::deque` is like vector, inserting to front is fast, other operations slower
- `std::stack` and `std::queue` adapt a vector/deque to their unique interface
- `std::set` is like Stanford Set. Main difference: `count()` instead of `contains()`
- `std::map` kinda like Stanford Map, but internally stores `pair<K, V>`
- `std::map`'s `operator[]` has auto-insertion, but `at()` does not
- `set/map` require a comparison operator (since they are internally sorted)
- you should practice using these containers yourself
- you should read the STL documentation for more details



Next time


Iterators



Homework

Finish GraphViz

Start "Cracking the Coding Interview in C++"



After Class Optional Example

CS 106B map/set examples, using the STL map/set